

DDL Commands

```
Create Table USERS (  
    user_id VARCHAR(100) Primary key,  
    username VARCHAR(100),  
    password VARCHAR(255),  
    user_email VARCHAR(100),  
    user_age INT  
);
```

```
Create table PARTS (  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    part_name VARCHAR(100),  
    part_png VARCHAR(255),  
    part_dimensions VARCHAR(255),  
    PRIMARY KEY(part_id, part_color)  
);
```

```
Create Table BUILDS (  
    build_id VARCHAR(100) Primary Key,  
    build_name VARCHAR(100),  
    build_png VARCHAR(100),  
    build_link VARCHAR(255),  
    build_age_rating INTEGER,  
    build_rating DECIMAL(4, 2),  
    build_release_year VARCHAR(10)  
);
```

```
Create Table THEMES (  
    theme_id VARCHAR(100) Primary Key,  
    theme_name VARCHAR(100),  
    theme_description VARCHAR(1000),  
    popular_build_id_1 VARCHAR(100),  
    popular_build_id_2 VARCHAR(100),  
    popular_build_id_3 VARCHAR(100),  
    FOREIGN KEY (popular_build_id_1) REFERENCES  
    BUILDS(build_id),
```

```
        FOREIGN KEY (popular_build_id_2) REFERENCES  
BUILDS(build_id),  
        FOREIGN KEY (popular_build_id_3) REFERENCES BUILDS(build_id)  
    );
```

```
Create Table SUPPLIERS (  
    supplier_id VARCHAR(100) Primary Key,  
    supplier_name VARCHAR(100),  
    supplier_region VARCHAR(100)  
);
```

```
Create Table INVENTORY (  
    user_id VARCHAR(100),  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    part_quantity INTEGER,  
    PRIMARY KEY(user_id, part_id, part_color),  
    FOREIGN KEY (user_id) REFERENCES USERS(user_id),  
    FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,  
part_color)  
);
```

```
Create Table BUILD_DETAILS (  
    build_id VARCHAR(100),  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    part_quantity INTEGER,  
    PRIMARY KEY(build_id, part_id, part_color),  
    FOREIGN KEY (build_id) REFERENCES BUILDS(build_id),  
    FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,  
part_color)  
);
```

```
Create Table FAVORITES(  
    user_id VARCHAR(100),  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    PRIMARY KEY(user_id, part_id, part_color),  
    FOREIGN KEY (user_id) REFERENCES USERS(user_id),
```

```
        FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,  
part_color)  
    );
```

```
Create Table REVIEWS(  
    user_id VARCHAR(100),  
    build_id VARCHAR(100),  
    review_text VARCHAR(1000),  
    PRIMARY KEY(user_id),  
    FOREIGN KEY (user_id) REFERENCES USERS(user_id),  
    FOREIGN KEY (build_id) REFERENCES BUILDS(build_id)  
);
```

```
Create Table BUILD_PRICING (  
    supplier_id VARCHAR(100),  
    build_id VARCHAR(100),  
    PRIMARY KEY(supplier_id, build_id),  
    FOREIGN KEY (build_id) REFERENCES BUILDS(build_id),  
    FOREIGN KEY (supplier_id) REFERENCES SUPPLIERS(supplier_id),  
    build_price DECIMAL(6, 2)  
);
```

```
Create Table PART_PRICING (  
    supplier_id VARCHAR(100),  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    PRIMARY KEY(supplier_id, part_id, part_color),  
    FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,  
part_color),  
    FOREIGN KEY (supplier_id) REFERENCES SUPPLIERS(supplier_id),  
    part_price DECIMAL(6, 2)  
);
```

```
Create Table BUILD_HAS_THEME (  
    theme_id VARCHAR(100),  
    build_id VARCHAR(100),  
    PRIMARY KEY(theme_id, build_id),  
    FOREIGN KEY (theme_id) REFERENCES THEMES(theme_id),  
    FOREIGN KEY (build_id) REFERENCES BUILDS(build_id)  
);
```

GET the Fields of each Table

```
mysql> show tables;
```

```
+-----+
| Tables_in_legoLab |
+-----+
| BUILDS             |
| BUILD_DETAILS      |
| BUILD_HAS_THEME    |
| BUILD_PRICING      |
| FAVORITES          |
| INVENTORY          |
| PARTS              |
| PART_PRICING       |
| REVIEWS            |
| SUPPLIERS          |
| THEMES             |
| USERS              |
+-----+
```

```
12 rows in set (0.00 sec)
```

```
mysql> describe builds;
```

```
+-----+-----+-----+-----+-----+-----+
| Field          | Type          | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| build_id       | varchar(100)  | NO   | PRI | NULL     |       |
| build_name     | varchar(100)  | YES  |     | NULL     |       |
| build_png      | varchar(100)  | YES  |     | NULL     |       |
| build_link     | varchar(255)  | YES  |     | NULL     |       |
| build_age_rating | int           | YES  |     | NULL     |       |
| build_rating   | decimal(4,2)  | YES  |     | NULL     |       |
| build_release_year | varchar(10)  | YES  |     | NULL     |       |
+-----+-----+-----+-----+-----+-----+
```

```
7 rows in set (0.00 sec)
```

```
mysql> describe BUILD_DETAILS;
```

```
+-----+-----+-----+-----+-----+-----+
| Field          | Type          | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| user_id       | varchar(100)  | NO   | PRI | NULL     |       |
| build_id      | varchar(100)  | NO   | PRI | NULL     |       |
| part_quantity | int           | YES  |     | NULL     |       |
+-----+-----+-----+-----+-----+-----+
```

```
3 rows in set (0.00 sec)
```

```
mysql> describe BUILD_HAS_THEME;
```

```
+-----+-----+-----+-----+-----+-----+
| Field      | Type          | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| theme_id   | varchar(100)  | NO   | PRI | NULL     |       |
| build_id   | varchar(100)  | NO   | PRI | NULL     |       |
+-----+-----+-----+-----+-----+-----+
```

```
2 rows in set (0.00 sec)
```

```
mysql> describe inventory;
```

```
+-----+-----+-----+-----+-----+-----+
| Field          | Type          | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| user_id       | varchar(100)  | NO   | PRI | NULL     |       |
| part_id       | varchar(100)  | NO   | PRI | NULL     |       |
| part_color    | varchar(100)  | NO   | PRI | NULL     |       |
| part_quantity | int           | YES  |     | NULL     |       |
+-----+-----+-----+-----+-----+-----+
```

```
4 rows in set (0.00 sec)
```

```
mysql> describe BUILD_PRICING;
```

Field	Type	Null	Key	Default	Extra
supplier_id	varchar(100)	NO	PRI	NULL	
build_id	varchar(100)	NO	PRI	NULL	
build_price	decimal(6,2)	YES		NULL	

3 rows in set (0.01 sec)

```
mysql> describe FAVORITES;
```

Field	Type	Null	Key	Default	Extra
user_id	varchar(100)	NO	PRI	NULL	
part_id	varchar(100)	NO	PRI	NULL	
part_color	varchar(100)	NO	PRI	NULL	

3 rows in set (0.01 sec)

```
mysql> describe PARTS;
```

Field	Type	Null	Key	Default	Extra
part_id	varchar(100)	NO	PRI	NULL	
part_color	varchar(100)	NO	PRI	NULL	
part_name	varchar(100)	YES		NULL	
part_png	varchar(255)	YES		NULL	
part_dimensions	varchar(255)	YES		NULL	

5 rows in set (0.00 sec)

```
mysql> describe PART_PRICING;
```

Field	Type	Null	Key	Default	Extra
supplier_id	varchar(100)	NO	PRI	NULL	
part_id	varchar(100)	NO	PRI	NULL	
part_color	varchar(100)	NO	PRI	NULL	
part_price	decimal(6,2)	YES		NULL	

4 rows in set (0.00 sec)

```
mysql> describe REVIEWS;
```

Field	Type	Null	Key	Default	Extra
user_id	varchar(100)	NO	PRI	NULL	
build_id	varchar(100)	YES	MUL	NULL	
review_text	varchar(1000)	YES		NULL	

3 rows in set (0.01 sec)

```
mysql> describe SUPPLIERS;
```

Field	Type	Null	Key	Default	Extra
supplier_id	varchar(100)	NO	PRI	NULL	
supplier_name	varchar(100)	YES		NULL	
supplier_region	varchar(100)	YES		NULL	

3 rows in set (0.01 sec)

```
mysql> describe THEMES;
```

Field	Type	Null	Key	Default	Extra
theme_id	varchar(100)	NO	PRI	NULL	
theme_name	varchar(100)	YES		NULL	
theme_description	varchar(1000)	YES		NULL	
popular_build_id_1	varchar(100)	YES	MUL	NULL	
popular_build_id_2	varchar(100)	YES	MUL	NULL	
popular_build_id_3	varchar(100)	YES	MUL	NULL	

6 rows in set (0.01 sec)

```
mysql> describe USERS;
```

Field	Type	Null	Key	Default	Extra
user_id	varchar(100)	NO	PRI	NULL	
username	varchar(100)	YES		NULL	
password	varchar(100)	YES		NULL	
user_email	varchar(100)	YES		NULL	
user_age	int	YES		NULL	

5 rows in set (0.01 sec)

Count of Each Table

```
mysql> select COUNT(*) from PARTS;
+-----+
| COUNT(*) |
+-----+
|      77948 |
+-----+
1 row in set (0.01 sec)
```

```
mysql> select COUNT(*) from BUILDS;
+-----+
| COUNT(*) |
+-----+
|      24195 |
+-----+
1 row in set (0.02 sec)
```

```
mysql> SELECT COUNT(*) FROM BUILD_HAS_THEME;
+-----+
| COUNT(*) |
+-----+
|      24198 |
+-----+
1 row in set (0.64 sec)
```

```
mysql> select COUNT(*) from BUILD_DETAILS;
+-----+
| COUNT(*) |
+-----+
|     857035 |
+-----+
1 row in set (2.07 sec)
```

```
mysql> select COUNT(*) from THEMES;
+-----+
| COUNT(*) |
+-----+
|      1046 |
+-----+
1 row in set (0.00 sec)
```

```
mysql> select COUNT(*) from PART_PRICING;
+-----+
| COUNT(*) |
+-----+
|          0 |
+-----+
1 row in set (0.00 sec)
```

```
mysql> SELECT COUNT(*) FROM SUPPLIERS;
+-----+
| COUNT(*) |
+-----+
|          0 |
+-----+
1 row in set (0.04 sec)
```

```
mysql> select COUNT(*) from REVIEWS;
+-----+
| COUNT(*) |
+-----+
|          0 |
+-----+
1 row in set (0.01 sec)
```

```
mysql> select COUNT(*) from USERS;
+-----+
| COUNT(*) |
+-----+
|         27 |
+-----+
1 row in set (0.00 sec)
```

```
mysql> select COUNT(*) from BUILD_PRICING;
+-----+
| COUNT(*) |
+-----+
|          0 |
+-----+
1 row in set (0.00 sec)
```

```
mysql> select COUNT(*) from FAVORITES;
+-----+
| COUNT(*) |
+-----+
|          0 |
+-----+
1 row in set (0.01 sec)
```

```
mysql> select COUNT(*) from INVENTORY;
+-----+
| COUNT(*) |
+-----+
|          9 |
+-----+
1 row in set (0.00 sec)
```

Queries:

1. Query: Find the Most Frequently Used Parts Across All Builds

This query identifies the parts most frequently used across all builds by joining the PARTS and Build_Details tables, aggregating by part_id and part_color.

SQL Concepts: **Join multiple relations, Aggregation with GROUP BY**

```
SELECT DISTINCT lp.part_id, lp.part_name, SUM(bd.part_quantity) AS
total_quantity_used
FROM PARTS AS lp
JOIN BUILD_DETAILS AS bd ON lp.part_id = bd.part_id
GROUP BY lp.part_id, lp.part_name
ORDER BY total_quantity_used DESC, lp.part_name
LIMIT 15;
```

```
mysql> SELECT DISTINCT lp.part_id, lp.part_name, SUM(bd.part_quantity) AS total_quantity_used
-> FROM PARTS AS lp
-> JOIN BUILD_DETAILS AS bd ON lp.part_id = bd.part_id
-> GROUP BY lp.part_id, lp.part_name
-> ORDER BY total_quantity_used DESC, lp.part_name
-> LIMIT 15;
```

part_id	part_name	total_quantity_used
3023	Plate 1 x 2	7006674
3004	Brick 1 x 2	5159880
3005	Brick 1 x 1	4457880
6141	Plate Round 1 x 1 with Solid Stud	3950080
3024	Plate 1 x 1	3946272
3003	Brick 2 x 2	2527254
3069b	Tile 1 x 2 with Groove	2459392
3710	Plate 1 x 4	2415899
3010	Brick 1 x 4	2182880
3020	Plate 2 x 4	2098440
3001	Brick 2 x 4	2094064
98138	Tile Round 1 x 1	1842749
3022	Plate 2 x 2	1675879
54200	Slope 30° 1 x 1 x 2/3 (Cheese Slope)	1378089
3021	Plate 2 x 3	1350404

15 rows in set (8 min 21.16 sec)

2. Query: Find Builds with the Largest Variety of Unique Parts

This query finds builds with the highest count of unique parts (based on part_id and part_color) used in each build, helping to identify builds that are the most complex in terms of part diversity.

SQL Concepts: **Join multiple relations, Aggregation with GROUP BY**

```
SELECT lb.build_id, lb.build_name, COUNT(DISTINCT bd.part_id, bd.part_color) AS
unique_part_count
FROM BUILDS AS lb
JOIN BUILD_DETAILS AS bd ON lb.build_id = bd.build_id
GROUP BY lb.build_id, lb.build_name
ORDER BY unique_part_count DESC, lb.build_name
LIMIT 15;
```

```
mysql> SELECT lb.build_id, lb.build_name, COUNT(DISTINCT bd.part_id, bd.part_color) AS unique_part_count
-> FROM BUILDS AS lb
-> JOIN BUILD_DETAILS AS bd ON lb.build_id = bd.build_id
-> GROUP BY lb.build_id, lb.build_name
-> ORDER BY unique_part_count DESC, lb.build_name
-> LIMIT 15;
```

build_id	build_name	unique_part_count
LEGO-Modulex-1	Unused Modulex parts sold by LEGO	1850
BIGBOX-1	The Ultimate Battle for Chima	1554
71741-1	NINJAGO City Gardens	1133
75978-1	Diagon Alley	1013
70620-1	NINJAGO City	915
21330-1	Home Alone	740
75331-1	The Razor Crest	723
10294-1	Titanic	692
70922-1	The Joker Manor	681
75192-1	Millennium Falcon	674
75313-1	AT-AT	658
10255-1	Assembly Square	649
70840-1	Welcome to Apocalypseburg!	648
70657-1	NINJAGO City Docks	640
10297-1	Boutique Hotel	635

15 rows in set (3.44 sec)

3. Query: Find Builds Released in the Last 5 Years with Above-Average Ratings

This query identifies builds released within the past five years that have a rating above the overall average for that time period.

SQL Concepts: **Subqueries, Aggregation with GROUP BY**

```
SELECT build_id, build_name, build_release_year, build_rating
FROM BUILDS
WHERE build_release_year >= 2019
AND build_rating > (
    SELECT AVG(build_rating)
    FROM BUILDS
    WHERE build_release_year >= 2019
)
```

ORDER BY build_release_year DESC, build_rating DESC
LIMIT 15;

```
mysql> SELECT build_id, build_name, build_release_year, build_rating
-> FROM BUILDS
-> WHERE build_release_year >= 2019
-> AND build_rating > (
->     SELECT AVG(build_rating)
->     FROM BUILDS
->     WHERE build_release_year >= 2019
-> )
-> ORDER BY build_release_year DESC, build_rating DESC
-> LIMIT 15;
```

build_id	build_name	build_release_year	build_rating
9781837250622-1	ReBuild Activity Cards: Animals	2025	5.00
53434-1	Spaceman Red Pencil Case Pop Up	2025	5.00
53500-1	Space Bus Molded Pencil Case	2025	4.90
77002-1	Cyclone vs. Metal Sonic	2025	4.90
9780241727416-1	LEGO Ideas Activity Book: Animals	2025	4.60
77053-1	Stargazing with Celeste	2025	4.20
9780241740859-1	Our Amazing Universe: Fantastic Building Ideas and Facts About Our Universe	2025	4.10
9781837250639-1	ReBuild Activity Cards: Magic	2025	4.00
9780794453343-1	Build and Play! Easter	2025	3.80
72035-1	Mario Kart â?? Toad's Garage	2025	3.70
910043-1	Forest Stronghold	2025	3.50
53389-1	Spaceman Blue Pencil Case Pop Up	2025	3.50
910042-1	Lost City	2025	3.30
77001-1	Sonicâ??s Campfire Clash	2025	3.00
53325-1	Ninjago Kai Molded Pencil Case	2025	2.60

15 rows in set (0.02 sec)

4. Find Builds with the Highest Number of Parts Used

This query lists the builds with the most parts used, based on the sum of part_quantity in Build_Details. This can be useful for identifying complex builds.

SQL Concepts: **Join multiple relations, Aggregation with GROUP BY**

```
SELECT lb.build_id, lb.build_name, SUM(bd.part_quantity) AS total_parts_used
FROM BUILDS AS lb
JOIN BUILD_DETAILS AS bd ON lb.build_id = bd.build_id
GROUP BY lb.build_id, lb.build_name
ORDER BY total_parts_used DESC
LIMIT 15;
```

```
mysql> SELECT lb.build_id, lb.build_name, SUM(bd.part_quantity) AS total_parts_used
-> FROM BUILDS AS lb
-> JOIN BUILD_DETAILS AS bd ON lb.build_id = bd.build_id
-> GROUP BY lb.build_id, lb.build_name
-> ORDER BY total_parts_used DESC
-> LIMIT 15;
```

build_id	build_name	total_parts_used
31203-1	World Map	11695
BIGBOX-1	The Ultimate Battle for Chima	9358
10294-1	Titanic	8636
10276-1	Colosseum	8389
75192-1	Millennium Falcon	6975
75313-1	AT-AT	5946
75331-1	The Razor Crest	5837
10256-1	Taj Mahal	5742
10189-1	Taj Mahal	5613
10299-1	Real Madrid â?? Santiago Bernab�u Stadium	5562
10284-1	Camp Nou - FC Barcelona	5476
SWMP-1	Star Wars / M&M Mosaic - Promo Set	5462
71043-1	Hogwarts Castle	5443
71741-1	NINJAGO City Gardens	5306
75978-1	Diagon Alley	5015

15 rows in set (3.25 sec)

INDEXING:

Query 1

Default Index

- Cost: 99847 for table scan
- Time: 23292 on table scan

```
-----+
| -> Limit: 15 row(s) (actual time=540151..540151 rows=15 loops=1)
-> Sort: total_quantity_used DESC, lp.part_name, limit input to 15 row(s) per chunk (actual time=540151..540151 rows=15 loops=1)
-> Table scan on <temporary> (actual time=540113..540132 rows=31944 loops=1)
-> Aggregate using temporary table (actual time=540113..540113 rows=31944 loops=1)
-> Nested loop inner join (cost=593901 rows=1.94e+6) (actual time=1.85..72274 rows=22.7e+6 loops=1)
-> Table scan on bd (cost=99847 rows=975033) (actual time=1.81..23292 rows=857035 loops=1)
-> Index lookup on lp using PRIMARY (part_id=bd.part_id) (cost=0.308 rows=1.99) (actual time=0.00822..0.0521 rows=26.5 loop
s=857035)
|
+-----+
```

Index 1:

Index on Parts.part_name

Why: This index will enhance the performance of queries that join the Parts table with the BUILD_DETAILS table, especially when retrieving part quantities.

CREATE INDEX idx_part_details ON PARTS (part_name);

```

-----+
| -> Limit: 15 row(s) (actual time=1.63e+6..1.63e+6 rows=15 loops=1)
|   -> Sort: total_quantity_used DESC, lp.part_name, limit input to 15 row(s) per chunk (actual time=1.63e+6..1.63e+6 rows=15 loops=1)
|     -> Table scan on <temporary> (actual time=1.63e+6..1.63e+6 rows=31944 loops=1)
|       -> Aggregate using temporary table (actual time=1.63e+6..1.63e+6 rows=31944 loops=1)
|         -> Nested loop inner join (cost=1.24e+6 rows=3.52e+6) (actual time=0.146..624243 rows=22.7e+6 loops=1)
|           -> Covering index scan on lp using idx_part_details (cost=10001 rows=83010) (actual time=0.0894..541 rows=77948 loops=1)
|             -> Index lookup on bd using part_id (part_id=lp.part_id) (cost=10.6 rows=42.4) (actual time=0.0647..7.95 rows=291 loops=779
48)
|
+-----+

```

Findings and Explanation: The actual execution time decreased from 23292 seconds to 541 seconds, showing a measurable improvement. The query cost decreased considerably to 10001 from 99847. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 2:

Index on Build_Details.part_quantity

Why: This index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when counting parts quantities across builds

CREATE INDEX idx_build_details ON Build_Details (part_quantity);

```

-----+
| -> Limit: 15 row(s) (actual time=486505..486505 rows=15 loops=1)
|   -> Sort: total_quantity_used DESC, lp.part_name, limit input to 15 row(s) per chunk (actual time=486504..486505 rows=15 loops=1)
|     -> Table scan on <temporary> (actual time=486465..486485 rows=31944 loops=1)
|       -> Aggregate using temporary table (actual time=486465..486465 rows=31944 loops=1)
|         -> Nested loop inner join (cost=1.27e+6 rows=1.94e+6) (actual time=3.51..55960 rows=22.7e+6 loops=1)
|           -> Covering index scan on bd using idx_build_details (cost=99874 rows=975033) (actual time=0.0581..3546 rows=857035 loops=1)
|             -> Index lookup on lp using PRIMARY (part_id=bd.part_id) (cost=0.999 rows=1.99) (actual time=0.0102..0.056 rows=26.5 loops=857035)
|
+-----+

```

Findings and Explanation: The actual execution time increased from 23292 seconds to 3546 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 3:

Index on Parts.part_name and Parts.part_dimensions

Why: This composite index will enhance the performance of queries that join the Build_Details table with the Part table, especially when retrieving part quantities and parts dimensions builds

CREATE INDEX idx_part_dim_col ON PARTS(part_name, part_dimensions);

```

-----+
| -> Limit: 15 row(s) (actual time=472192..472192 rows=15 loops=1)
  -> Sort: total_quantity_used DESC, lp.part_name, limit input to 15 row(s) per chunk (actual time=472192..472192 rows=15 loops=1)
    -> Table scan on <temporary> (actual time=472152..472172 rows=31944 loops=1)
      -> Aggregate using temporary table (actual time=472152..472152 rows=31944 loops=1)
        -> Nested loop inner join (cost=539734 rows=1.94e+6) (actual time=1.82..56658 rows=22.7e+6 loops=1)
          -> Table scan on bd (cost=101475 rows=975033) (actual time=1.77..7036 rows=857035 loops=1)
            -> Index lookup on lp using PRIMARY (part_id=bd.part_id) (cost=0.251 rows=1.99) (actual time=0.0084..0.0529 rows=26.5 loops=1)
          =857035)
        |
      +-----+

```

Findings and Explanation: The actual execution time decreased from 23292 seconds to 7036 seconds, showing a measurable improvement. The query cost increased from 98860 to 99847. The Indexing strategy did not cause a significant improvement.

Query 2

Default Index

- Cost: 2428 on scan
- Time: 3183 on limit

```

-----+
| -> Limit: 15 row(s) (actual time=3183..3183 rows=15 loops=1)
  -> Sort: unique_part_count DESC, lb.build_name, limit input to 15 row(s) per chunk (actual time=3183..3183 rows=15 loops=1)
    -> Stream results (cost=308684 rows=284827) (actual time=24.4..3168 rows=15770 loops=1)
      -> Group aggregate: count(distinct bd.part_id, bd.part_color) (cost=308684 rows=284827) (actual time=24.4..3145 rows=15770 loops=1)
        -> Nested loop inner join (cost=193698 rows=1.15e+6) (actual time=24.3..1339 rows=857035 loops=1)
          -> Sort: lb.build_id, lb.build_name (cost=2428 rows=23559) (actual time=24.3..33.8 rows=24195 loops=1)
            -> Table scan on lb (cost=2428 rows=23559) (actual time=0.111..8.99 rows=24195 loops=1)
              -> Covering index lookup on bd using PRIMARY (build_id=lb.build_id) (cost=3.24 rows=48.8) (actual time=0.0118..0.0492 rows=35.4 loops=24195)
            |
          +-----+

```

Index 1:

Index on Build_Details.part_quantity

Why: This index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when retrieving part quantities.

CREATE INDEX idx_part_details ON Build_Details (part_quantity);

```

-----+
| -> Limit: 15 row(s) (actual time=3128..3128 rows=15 loops=1)
  -> Sort: unique_part_count DESC, lb.build_name, limit input to 15 row(s) per chunk (actual time=3128..3128 rows=15 loops=1)
    -> Stream results (cost=308711 rows=284827) (actual time=43.3..3114 rows=15770 loops=1)
      -> Group aggregate: count(distinct bd.part_id, bd.part_color) (cost=308711 rows=284827) (actual time=43.3..3091 rows=15770 loops=1)
        -> Nested loop inner join (cost=193725 rows=1.15e+6) (actual time=43.3..1333 rows=857035 loops=1)
          -> Sort: lb.build_id, lb.build_name (cost=2455 rows=23559) (actual time=43.2..52.8 rows=24195 loops=1)
            -> Table scan on lb (cost=2455 rows=23559) (actual time=3.74..26.4 rows=24195 loops=1)
              -> Covering index lookup on bd using PRIMARY (build_id=lb.build_id) (cost=3.24 rows=48.8) (actual time=0.0113..0.0483 rows=35.4 loops=24195)
            |
          +-----+

```

Findings and Explanation: The actual execution time decreased from 3183 seconds to 26.4 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed

of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 2:

Index on Builds.release_year

Why: This index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when counting unique release across builds

CREATE INDEX idx_build_year ON BUILDS (build_release_year);

```
-----+
| -> Limit: 15 row(s) (actual time=3095..3095 rows=15 loops=1)
|   -> Sort: unique_part_count DESC, lb.build_name, limit input to 15 row(s) per chunk (actual time=3095..3095 rows=15 loops=1)
|     -> Stream results (cost=308684 rows=284827) (actual time=17.6..3080 rows=15770 loops=1)
|       -> Group aggregate: count(distinct bd.part_id, bd.part_color) (cost=308684 rows=284827) (actual time=17.6..3057 rows=15770 loops=1)
|         -> Nested loop inner join (cost=193698 rows=1.15e+6) (actual time=17.6..1319 rows=857035 loops=1)
|           -> Sort: lb.build_id, lb.build_name (cost=2428 rows=23559) (actual time=17.5..28.1 rows=24195 loops=1)
|             -> Table scan on lb (cost=2428 rows=23559) (actual time=0.0371..5.89 rows=24195 loops=1)
|               -> Covering index lookup on bd using PRIMARY (build_id=lb.build_id) (cost=3.24 rows=48.8) (actual time=0.0118..0.0488 rows=
35.4 loops=24195)
|
|-----+
```

Findings and Explanation: The actual execution time decreased from 3183 seconds to 5.89 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 3:

Index on BUILDS.age_rating

Why: This composite index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when retrieving part quantities and counting unique parts across builds

CREATE INDEX idx_build_age ON BUILDS(build_age_rating);

```
-----+
| -> Limit: 15 row(s) (actual time=3376..3376 rows=15 loops=1)
|   -> Sort: unique_part_count DESC, lb.build_name, limit input to 15 row(s) per chunk (actual time=3376..3376 rows=15 loops=1)
|     -> Stream results (cost=308684 rows=284827) (actual time=18.3..3360 rows=15770 loops=1)
|       -> Group aggregate: count(distinct bd.part_id, bd.part_color) (cost=308684 rows=284827) (actual time=18.3..3334 rows=15770 loops=1)
|         -> Nested loop inner join (cost=193698 rows=1.15e+6) (actual time=18.3..1435 rows=857035 loops=1)
|           -> Sort: lb.build_id, lb.build_name (cost=2428 rows=23559) (actual time=18.2..28.3 rows=24195 loops=1)
|             -> Table scan on lb (cost=2428 rows=23559) (actual time=0.0467..5.96 rows=24195 loops=1)
|               -> Covering index lookup on bd using PRIMARY (build_id=lb.build_id) (cost=3.24 rows=48.8) (actual time=0.0125..0.0532 rows=
35.4 loops=24195)
|
|-----+
```

Findings and Explanation: The actual execution time decreased from 3183 seconds to 5.96 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Query 3

Default Index

- Cost: 398.13 on limit
- Time: 7.458 on limit

```
mysql> EXPLAIN ANALYZE SELECT build_id, build_name, build_release_year, build_rating FROM small_BUILDS WHERE build_release_year >= 2019 AND build_rating > (
  SELECT AVG(build_rating) FROM small_BUILDS WHERE build_release_year >= 2019 ) ORDER BY build_release_year DESC, build_rating DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (cost=398.13 rows=15) (actual time=7.450..7.458 rows=15 loops=1)
   -> Sort: small_BUILDS.build_release_year DESC, small_BUILDS.build_rating DESC, limit input to 15 row(s) per chunk (cost=398.13 rows=4807) (actual time=7.450..7.458 rows=15 loops=1)
   -> Filter: ((small_BUILDS.build_release_year >= 2019) and (small_BUILDS.build_rating > (select #2))) (cost=398.13 rows=4807) (actual time=3.610..7.244 rows=15 loops=1)
   -> Table scan on small_BUILDS (cost=398.13 rows=4807) (actual time=0.037..3.039 rows=5000 loops=1)
   -> Select #2 (subquery in condition; run only once)
       -> Aggregate: avg(small_BUILDS.build_rating) (cost=665.17 rows=1) (actual time=3.525..3.526 rows=1 loops=1)
       -> Filter: (small_BUILDS.build_release_year >= 2019) (cost=504.95 rows=1602) (actual time=0.006..3.344 rows=1082 loops=1)
       -> Table scan on small_BUILDS (cost=504.95 rows=4807) (actual time=0.004..2.862 rows=5000 loops=1)
+-----+
| (END) |
+-----+
```

Index 1:

Index on Builds.builds_release_year

Why: This index will help improve the speed of the query when we are checking by indexing through build_release_year.

CREATE INDEX idx_build_year ON small_BUILDS(build_release_year);

```
mysql> EXPLAIN ANALYZE SELECT build_id, build_name, build_release_year, build_rating FROM small_BUILDS WHERE build_release_year >= 2019 AND build_rating > (
  SELECT AVG(build_rating) FROM small_BUILDS WHERE build_release_year >= 2019 ) ORDER BY build_release_year DESC, build_rating DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (cost=398.13 rows=15) (actual time=7.404..7.407 rows=15 loops=1)
   -> Sort: small_BUILDS.build_release_year DESC, small_BUILDS.build_rating DESC, limit input to 15 row(s) per chunk (cost=398.13 rows=4807) (actual time=7.404..7.407 rows=15 loops=1)
   -> Filter: ((small_BUILDS.build_release_year >= 2019) and (small_BUILDS.build_rating > (select #2))) (cost=398.13 rows=4807) (actual time=3.654..7.204 rows=15 loops=1)
   -> Table scan on small_BUILDS (cost=398.13 rows=4807) (actual time=0.081..3.015 rows=5000 loops=1)
   -> Select #2 (subquery in condition; run only once)
       -> Aggregate: avg(small_BUILDS.build_rating) (cost=665.17 rows=1) (actual time=3.524..3.525 rows=1 loops=1)
       -> Filter: (small_BUILDS.build_release_year >= 2019) (cost=504.95 rows=1602) (actual time=0.006..3.386 rows=1082 loops=1)
       -> Table scan on small_BUILDS (cost=504.95 rows=4807) (actual time=0.004..2.857 rows=5000 loops=1)
+-----+
| 1 row in set, 2 warnings (0.01 sec) |
+-----+
```

Findings and Explanation: The actual execution time did not decrease at all. The query cost remained roughly the same. The Indexing strategy did not cause a significant improvement.

Index 2:

Index on BUILDS.build_rating

Why: An index on these columns will speed up queries that filter on build_rating, like the one that identifies builds released above a certain rating.

CREATE INDEX idx_build_rating ON BUILDS(build_rating);

```
mysql> EXPLAIN ANALYZE SELECT build_id, build_name, build_release_year, build_rating FROM small_BUILDS WHERE build_release_year >= 2019 AND build_rating > (
  SELECT AVG(build_rating) FROM small_BUILDS WHERE build_release_year >= 2019 ) ORDER BY build_release_year DESC, build_rating DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (cost=504.95 rows=15) (actual time=4.017..4.019 rows=15 loops=1)
  -> Sort: small_BUILDS.build_release_year DESC, small_BUILDS.build_rating DESC, limit input to 15 row(s) per chunk (cost=504.95 rows=4807) (actual time=4.017..4.019 rows=15 loops=1)
    -> Filter: ((small_BUILDS.build_release_year >= 2019) and (small_BUILDS.build_rating > (select #2))) (cost=504.95 rows=4807) (actual time=0.054..3.811 rows=15 loops=1)
      -> Table scan on small_BUILDS (cost=504.95 rows=4807) (actual time=0.024..3.185 rows=5000 loops=1)
        -> Select #2 (subquery in condition; run only once)
          -> Aggregate: avg(small_BUILDS.build_rating) (cost=665.17 rows=1) (actual time=3.586..3.586 rows=1 loops=1)
            -> Filter: (small_BUILDS.build_release_year >= 2019) (cost=504.95 rows=1602) (actual time=0.049..3.449 rows=1082 loops=1)
              -> Table scan on small_BUILDS (cost=504.95 rows=4807) (actual time=0.037..2.886 rows=5000 loops=1)
+-----+
1 row in set, 2 warnings (0.01 sec)
```

Findings and Explanation: The actual execution time remained the same. The query cost actually increased. The indexing strategy reduced the speed of the query by decreasing the speed of joins, aggregations, and sorting operations, leading to slower execution times and less efficient resource usage.

Index 3:

Index on BUILDS.builds_release_year and BUILDS.builds_rating

Why: Since the Build_Themes table is often joined with the BUILDS table, indexing theme_id will speed up join operations when retrieving theme information related to builds.

CREATE INDEX idx_release_year_rating ON BUILDS (build_release_year, build_rating);

```
mysql> CREATE INDEX idx_release_year_rating ON small_BUILDS (build_release_year, build_rating);
Query OK, 0 rows affected (0.17 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> EXPLAIN ANALYZE SELECT build_id, build_name, build_release_year, build_rating FROM small_BUILDS WHERE build_release_year >= 2019 AND build_rating > (
  SELECT AVG(build_rating) FROM small_BUILDS WHERE build_release_year >= 2019 ) ORDER BY build_release_year DESC, build_rating DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (cost=2.56 rows=5) (actual time=0.060..0.133 rows=15 loops=1)
  -> Filter: ((small_BUILDS.build_release_year >= 2019) and (small_BUILDS.build_rating > (select #2))) (cost=2.56 rows=5) (actual time=0.059..0.130 rows=15 loops=1)
    -> Index scan on small_BUILDS using idx_release_year_rating (reverse) (cost=2.56 rows=29) (actual time=0.050..0.115 rows=15 loops=1)
      -> Select #2 (subquery in condition; run only once)
        -> Aggregate: avg(small_BUILDS.build_rating) (cost=665.17 rows=1) (actual time=5.381..5.381 rows=1 loops=1)
          -> Filter: (small_BUILDS.build_release_year >= 2019) (cost=504.95 rows=1602) (actual time=4.033..5.151 rows=1082 loops=1)
            -> Covering index scan on small_BUILDS using idx_release_year_rating (cost=504.95 rows=4807) (actual time=0.036..4.216 rows=5000 loops=1)
+-----+
1 row in set, 0 warnings (0.01 sec)
```

Findings and Explanation: The actual execution time decreased greatly from 7.45 seconds to 0.06 seconds, showing a measurable improvement. The query cost also decreased considerably to 2.56 from 398.13. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Query 4

Default Index

- Cost: 2428 on table scan
- Time: 24.8 on limit


```

-----+
| -> Limit: 15 row(s) (actual time=3680..3680 rows=15 loops=1)
|   -> Sort: total_parts_used DESC, limit input to 15 row(s) per chunk (actual time=3680..3680 rows=15 loops=1)
|     -> Table scan on <temporary> (actual time=3668..3675 rows=15770 loops=1)
|       -> Aggregate using temporary table (actual time=3668..3668 rows=15770 loops=1)
|         -> Nested loop inner join (cost=123591 rows=1.15e+6) (actual time=0.161..1553 rows=857035 loops=1)
|           -> Table scan on lb (cost=2428 rows=23559) (actual time=0.104..24.8 rows=24195 loops=1)
|             -> Index lookup on bd using PRIMARY (build_id=lb.build_id) (cost=0.262 rows=48.8) (actual time=0.0125..0.0577 rows=35.4 loops=1)
|
ps=24195
|
+-----+

```

Index 1:

Index on Build_Details.part_quantity

Why: This index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when retrieving part quantities.

CREATE INDEX idx_part_details ON Build_Details (part_quantity);

```

-----+
| -> Limit: 15 row(s) (actual time=3678..3678 rows=15 loops=1)
|   -> Sort: total_parts_used DESC, limit input to 15 row(s) per chunk (actual time=3678..3678 rows=15 loops=1)
|     -> Table scan on <temporary> (actual time=3666..3672 rows=15770 loops=1)
|       -> Aggregate using temporary table (actual time=3666..3666 rows=15770 loops=1)
|         -> Nested loop inner join (cost=123591 rows=1.15e+6) (actual time=0.148..1547 rows=857035 loops=1)
|           -> Table scan on lb (cost=2428 rows=23559) (actual time=0.0942..26.4 rows=24195 loops=1)
|             -> Index lookup on bd using PRIMARY (build_id=lb.build_id) (cost=0.262 rows=48.8) (actual time=0.0123..0.0573 rows=35.4 loops=1)
|
ps=24195
|
+-----+

```

Findings and Explanation: The actual execution time remained roughly the same. The query cost remained roughly the same. The indexing strategy did not optimize and the query decreased in speed with less efficient resource use.

Index 2:

Index on Builds.release_year

Why: This index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when counting unique release across builds

CREATE INDEX idx_build_year ON BUILDS (build_release_year);

```

-----+
| -> Limit: 15 row(s) (actual time=3719..3719 rows=15 loops=1)
|   -> Sort: total_parts_used DESC, limit input to 15 row(s) per chunk (actual time=3719..3719 rows=15 loops=1)
|     -> Table scan on <temporary> (actual time=3707..3713 rows=15770 loops=1)
|       -> Aggregate using temporary table (actual time=3707..3707 rows=15770 loops=1)
|         -> Nested loop inner join (cost=123591 rows=1.15e+6) (actual time=0.0727..1544 rows=857035 loops=1)
|           -> Table scan on lb (cost=2428 rows=23559) (actual time=0.046..25 rows=24195 loops=1)
|             -> Index lookup on bd using PRIMARY (build_id=lb.build_id) (cost=0.262 rows=48.8) (actual time=0.0125..0.0574 rows=35.4 loops=1)
|
ps=24195
|
+-----+

```

Findings and Explanation: The actual execution time remained roughly the same. The query cost remained roughly the same. The indexing strategy did not optimize and the query decreased in speed with less efficient resource use.

Index 3:

Index on BUILDS.age_rating

Why: This composite index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when retrieving part quantities and counting unique parts across builds

CREATE INDEX idx_build_age ON BUILDS(build_age_rating);

```
-----+
| -> Limit: 15 row(s) (actual time=3605..3605 rows=15 loops=1)
|   -> Sort: total_parts_used DESC, limit input to 15 row(s) per chunk (actual time=3605..3605 rows=15 loops=1)
|     -> Table scan on <temporary> (actual time=3593..3599 rows=15770 loops=1)
|       -> Aggregate using temporary table (actual time=3593..3593 rows=15770 loops=1)
|         -> Nested loop inner join (cost=123591 rows=1.15e+6) (actual time=0.147..1518 rows=857035 loops=1)
|           -> Table scan on lb (cost=2428 rows=23559) (actual time=0.102..23.8 rows=24195 loops=1)
|             -> Index lookup on bd using PRIMARY (build_id=lb.build_id) (cost=0.262 rows=48.8) (actual time=0.0121..0.0564 rows=35.4 loops=1)
|
ps=24195
|
+-----+
```

Findings and Explanation: The actual execution time is roughly the same. The query cost remained roughly the same. The indexing strategy did not optimize and the query decreased in speed with less efficient resource use.

Final Choice

After evaluating the various indexing configurations, we ultimately selected the index on part_quantity for BUILD_DETAILS table and combined index on build_year_release and build_rating for BUILDS table and index on part_name for PARTS table as our final design. This decision was based on the consistent performance gains observed across multiple queries, particularly in those that build_release_year and build_rating the most. The improvements were significant enough to justify the added overhead of maintaining the index, and we anticipate that it will enhance user experience by reducing query response times.