

## **DDL Commands**

```
Create Table USERS (  
    user_id VARCHAR(100) Primary key,  
    username VARCHAR(100),  
    password VARCHAR(100),  
    user_email VARCHAR(100),  
    user_age INT  
);
```

```
Create table PARTS (  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    part_name VARCHAR(100),  
    part_png VARCHAR(255),  
    part_dimensions VARCHAR(255),  
    PRIMARY KEY(part_id, part_color)  
);
```

```
Create Table BUILDS (  
    build_id VARCHAR(100) Primary Key,  
    build_name VARCHAR(100),  
    build_png VARCHAR(100),  
    build_link VARCHAR(255),  
    build_age_rating INTEGER,  
    build_rating DECIMAL(4, 2),  
    build_release_year VARCHAR(10)  
);
```

```
Create Table THEMES (  
    theme_id VARCHAR(100) Primary Key,  
    theme_name VARCHAR(100),  
    theme_description VARCHAR(1000),  
    popular_build_id_1 VARCHAR(100),  
    popular_build_id_2 VARCHAR(100),  
    popular_build_id_3 VARCHAR(100),  
    FOREIGN KEY (popular_build_id_1) REFERENCES  
    BUILDS(build_id),
```

```
        FOREIGN KEY (popular_build_id_2) REFERENCES  
BUILDS(build_id),  
        FOREIGN KEY (popular_build_id_3) REFERENCES BUILDS(build_id)  
    );
```

```
Create Table SUPPLIERS (  
    supplier_id VARCHAR(100) Primary Key,  
    supplier_name VARCHAR(100),  
    supplier_region VARCHAR(100)  
);
```

```
Create Table INVENTORY (  
    user_id VARCHAR(100),  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    part_quantity INTEGER,  
    PRIMARY KEY(user_id, part_id, part_color),  
    FOREIGN KEY (user_id) REFERENCES USERS(user_id),  
    FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,  
part_color)  
);
```

```
Create Table BUILD_DETAILS (  
    build_id VARCHAR(100),  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    part_quantity INTEGER,  
    PRIMARY KEY(build_id, part_id, part_color),  
    FOREIGN KEY (build_id) REFERENCES BUILDS(build_id),  
    FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,  
part_color)  
);
```

```
Create Table FAVORITES(  
    user_id VARCHAR(100),  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    PRIMARY KEY(user_id, part_id, part_color),  
    FOREIGN KEY (user_id) REFERENCES USERS(user_id),
```

```
        FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,
part_color)
    );
    LOAD DATA LOCAL INFILE 'storage-legolab/data/PARTS.csv' INTO TABLE
table_name FIELDS TERMINATED BY ';' ENCLOSED BY '"' LINES
TERMINATED BY '\r\n'; IGNORE 1 ROWS;
```

```
Create Table REVIEWS(
    user_id VARCHAR(100),
    build_id VARCHAR(100),
    review_text VARCHAR(1000),
    PRIMARY KEY(user_id),
    FOREIGN KEY (user_id) REFERENCES USERS(user_id),
    FOREIGN KEY (build_id) REFERENCES BUILDS(build_id)
);
```

```
Create Table BUILD_PRICING (
    supplier_id VARCHAR(100),
    build_id VARCHAR(100),
    PRIMARY KEY(supplier_id, build_id),
    FOREIGN KEY (build_id) REFERENCES BUILDS(build_id),
    FOREIGN KEY (supplier_id) REFERENCES SUPPLIERS(supplier_id),
    build_price DECIMAL(6, 2)
);
```

```
Create Table PART_PRICING (
    supplier_id VARCHAR(100),
    part_id VARCHAR(100),
    part_color VARCHAR(100),
    PRIMARY KEY(supplier_id, part_id, part_color),
    FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,
part_color),
    FOREIGN KEY (supplier_id) REFERENCES SUPPLIERS(supplier_id),
    part_price DECIMAL(6, 2)
);
```

```
Create Table BUILD_HAS_THEME (
    theme_id VARCHAR(100),
    build_id VARCHAR(100),
    PRIMARY KEY(theme_id, build_id),
```

```
FOREIGN KEY (theme_id) REFERENCES THEMES(theme_id),
FOREIGN KEY (build_id) REFERENCES BUILDS(build_id)
);
```

### GET the Fields of each Table

```
mysql> show tables;
+-----+
| Tables_in_legoLab |
+-----+
| BUILDS             |
| BUILD_DETAILS      |
| BUILD_HAS_THEME    |
| BUILD_PRICING      |
| FAVORITES          |
| INVENTORY          |
| PARTS              |
| PART_PRICING       |
| REVIEWS            |
| SUPPLIERS          |
| THEMES             |
| USERS              |
+-----+
12 rows in set (0.00 sec)
```

```
mysql> describe BUILD_DETAILS;
```

Field	Type	Null	Key	Default	Extra
user_id	varchar(100)	NO	PRI	NULL	
build_id	varchar(100)	NO	PRI	NULL	
part_quantity	int	YES		NULL	

```
3 rows in set (0.00 sec)
```

```
mysql> describe BUILD_HAS_THEME;
```

Field	Type	Null	Key	Default	Extra
theme_id	varchar(100)	NO	PRI	NULL	
build_id	varchar(100)	NO	PRI	NULL	

```
2 rows in set (0.00 sec)
```

```
mysql> describe BUILD_PRICING;
```

Field	Type	Null	Key	Default	Extra
supplier_id	varchar(100)	NO	PRI	NULL	
build_id	varchar(100)	NO	PRI	NULL	
build_price	decimal(6,2)	YES		NULL	

3 rows in set (0.01 sec)

```
mysql> describe FAVORITES;
```

Field	Type	Null	Key	Default	Extra
user_id	varchar(100)	NO	PRI	NULL	
part_id	varchar(100)	NO	PRI	NULL	
part_color	varchar(100)	NO	PRI	NULL	

3 rows in set (0.01 sec)

```
mysql> describe PARTS;
```

Field	Type	Null	Key	Default	Extra
part_id	varchar(100)	NO	PRI	NULL	
part_color	varchar(100)	NO	PRI	NULL	
part_name	varchar(100)	YES		NULL	
part_png	varchar(255)	YES		NULL	
part_dimensions	varchar(255)	YES		NULL	

5 rows in set (0.00 sec)

```
mysql> describe PART_PRICING;
```

Field	Type	Null	Key	Default	Extra
supplier_id	varchar(100)	NO	PRI	NULL	
part_id	varchar(100)	NO	PRI	NULL	
part_color	varchar(100)	NO	PRI	NULL	
part_price	decimal(6,2)	YES		NULL	

4 rows in set (0.00 sec)

```
mysql> describe REVIEWS;
```

Field	Type	Null	Key	Default	Extra
user_id	varchar(100)	NO	PRI	NULL	
build_id	varchar(100)	YES	MUL	NULL	
review_text	varchar(1000)	YES		NULL	

3 rows in set (0.01 sec)

```
mysql> describe SUPPLIERS;
```

Field	Type	Null	Key	Default	Extra
supplier_id	varchar(100)	NO	PRI	NULL	
supplier_name	varchar(100)	YES		NULL	
supplier_region	varchar(100)	YES		NULL	

3 rows in set (0.01 sec)

```
mysql> describe USERS;
```

Field	Type	Null	Key	Default	Extra
user_id	varchar(100)	NO	PRI	NULL	
username	varchar(100)	YES		NULL	
password	varchar(100)	YES		NULL	
user_email	varchar(100)	YES		NULL	
user_age	int	YES		NULL	

5 rows in set (0.01 sec)

```
mysql> describe THEMES;
```

Field	Type	Null	Key	Default	Extra
theme_id	varchar(100)	NO	PRI	NULL	
theme_name	varchar(100)	YES		NULL	
theme_description	varchar(1000)	YES		NULL	
popular_build_id_1	varchar(100)	YES	MUL	NULL	
popular_build_id_2	varchar(100)	YES	MUL	NULL	
popular_build_id_3	varchar(100)	YES	MUL	NULL	

6 rows in set (0.01 sec)

### Count of Each Table

```
mysql> SELECT COUNT(*) FROM PARTS;
+-----+
| COUNT(*) |
+-----+
|      77949 |
+-----+
1 row in set (10.25 sec)
```

```
mysql> SELECT COUNT(*) FROM BUILDS;
+-----+
| COUNT(*) |
+-----+
|      24196 |
+-----+
1 row in set (4.79 sec)
```

```
mysql> SELECT COUNT(*) FROM USERS;
+-----+
| COUNT(*) |
+-----+
|          0 |
+-----+
1 row in set (0.08 sec)
```

```
mysql> SELECT COUNT(*) FROM SUPPLIERS;
+-----+
| COUNT(*) |
+-----+
|          0 |
+-----+
1 row in set (0.04 sec)
```

Queries:

## 1. Query: Find the Most Frequently Used Parts Across All Builds

This query identifies the parts most frequently used across all builds by joining the PARTS and Build\_Details tables, aggregating by part\_id and part\_color.

SQL Concepts: **Join multiple relations, Aggregation with GROUP BY**

```
SELECT lp.part_id, lp.part_color, lp.part_name, SUM(bd.part_quantity) AS  
total_quantity_used  
FROM PARTS AS lp  
JOIN BUILD_DETAILS AS bd ON lp.part_id = bd.part_id  
GROUP BY lp.part_id, lp.part_color, lp.part_name  
ORDER BY total_quantity_used DESC, lp.part_name  
LIMIT 15;
```

## 2. Query: List Themes with the Highest Average Build Rating

This query calculates the average rating for each theme by joining Build\_Themes with BUILDS and grouping by theme. The output is ordered by average rating in descending order.

SQL Concepts: **Join multiple relations, Aggregation with GROUP BY**

```
SELECT bt.theme_name, AVG(lb.build_rating) AS avg_rating FROM Build_Themes AS  
bt JOIN BUILDS AS lb ON bt.theme_id = lb.build_id GROUP BY bt.theme_name  
ORDER BY avg_rating DESC, bt.theme_name LIMIT 15;
```

## 3. Query: Find Builds Released in the Last 5 Years with Above-Average Ratings

This query identifies builds released within the past five years that have a rating above the overall average for that time period.

SQL Concepts: **Subqueries, Aggregation with GROUP BY**

```
SELECT build_id, build_name, build_release_year, build_rating  
FROM BUILDS  
WHERE build_release_year >= 2019  
AND build_rating > (
```



```

SELECT AVG(build_rating)
FROM BUILDS
WHERE build_release_year >= 2019
)
ORDER BY build_release_year DESC, build_rating DESC
LIMIT 15;

```

```

mysql> SELECT build_id, build_name, build_release_year, build_rating FROM BUILDS WHERE build_release_year >= 2019 AND build_rating > (
g) FROM BUILDS WHERE build_release_year >= 2019 ) ORDER BY build_release_year DESC, build_rating DESC LIMIT 15;

```

build_id	build_name	build_release_year	build_rating
9781837250622-1	ReBuild Activity Cards: Animals	2025	5.00
53434-1	Spaceman Red Pencil Case Pop Up	2025	5.00
53500-1	Space Bus Molded Pencil Case	2025	4.90
77002-1	Cyclone vs. Metal Sonic	2025	4.90
9780241727416-1	LEGO Ideas Activity Book: Animals	2025	4.60
77053-1	Stargazing with Celeste	2025	4.20
9780241740959-1	Our Amazing Universe: Fantastic Building Ideas and Facts About Our Universe	2025	4.10
9781837250639-1	ReBuild Activity Cards: Magic	2025	4.00
9780794453343-1	Build and Play! Easter	2025	3.80
72035-1	Mario Kart - Toad's Garage	2025	3.70
910043-1	Forest Stronghold	2025	3.50
53389-1	Spaceman Blue Pencil Case Pop Up	2025	3.50
910042-1	Lost City	2025	3.30
77001-1	Sonic's Campfire Clash	2025	3.00
53325-1	Ninjago Kai Molded Pencil Case	2025	2.60

15 rows in set, 2 warnings (0.81 sec)

#### 4. Query: Find Builds with the Largest Variety of Unique Parts

This query finds builds with the highest count of unique parts (based on part\_id and part\_color) used in each build, helping to identify builds that are the most complex in terms of part diversity.

SQL Concepts: **Join multiple relations, Aggregation with GROUP BY**

```

SELECT lb.build_id, lb.build_name, COUNT(DISTINCT bd.part_id, bd.part_color) AS
unique_part_count
FROM BUILDS AS lb
JOIN Build_Details AS bd ON lb.build_id = bd.build_id
GROUP BY lb.build_id, lb.build_name
ORDER BY unique_part_count DESC, lb.build_name
LIMIT 15;

```

#### 5. Find Builds with the Highest Number of Parts Used

This query lists the builds with the most parts used, based on the sum of part\_quantity in Build\_Details. This can be useful for identifying complex builds.

SQL Concepts: **Join multiple relations, Aggregation with GROUP BY**

```

SELECT lb.build_id, lb.build_name, SUM(bd.part_quantity) AS total_parts_used

```

```
FROM BUILDS AS lb
JOIN Build_Details AS bd ON lb.build_id = bd.build_id
GROUP BY lb.build_id, lb.build_name
ORDER BY total_parts_used DESC
LIMIT 15;
```

## **INDEXING:**

### **Query 1**

Default Index

- Cost: 802.05
- Time: 7.836..7.840

### **Index 1:**

Index on Build\_Details.part\_id and Build\_Details.build\_id

Why: This composite index will enhance the performance of queries that join the Build\_Details table with the BUILDS table, especially when retrieving part quantities and counting unique parts across builds.

```
CREATE INDEX idx_part_build ON Build_Details (part_id, build_id);
```

Findings and Explanation: The actual execution time decreased from 6.296 seconds to 4.126 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

### **Index 2:**

Index on BUILDS.build\_release\_year and BUILDS.build\_rating

Why: An index on these columns will speed up queries that filter on build\_release\_year and those that check for build\_rating, like the one that identifies builds released within a certain timeframe or above a certain rating.

```
CREATE INDEX idx_release_year_rating ON BUILDS (build_release_year,
build_rating);
```

Findings and Explanation: The actual execution time decreased from 7.428 seconds to 6.136 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 3:

Index on Build\_Themes.theme\_id

Why: Since the Build\_Themes table is often joined with the BUILDS table, indexing theme\_id will speed up join operations when retrieving theme information related to builds.

```
CREATE INDEX idx_theme_id ON Build_Themes (theme_id);
```

Findings and Explanation: The actual execution time decreased from 6.695 seconds to 5.187 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

## Query 2

Default Index

- Cost: 562.05
- Time: 7.836..7.840

Index 1:

Index on Build\_Details.part\_id and Build\_Details.build\_id

Why: This composite index will enhance the performance of queries that join the Build\_Details table with the BUILDS table, especially when retrieving part quantities and counting unique parts across builds.

```
CREATE INDEX idx_part_build ON Build_Details (part_id, build_id);
```

Findings and Explanation: The actual execution time decreased from 6.296 seconds to 4.126 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed

of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 2:

Index on BUILDS.build\_release\_year and BUILDS.build\_rating

Why: An index on these columns will speed up queries that filter on build\_release\_year and those that check for build\_rating, like the one that identifies builds released within a certain timeframe or above a certain rating.

```
CREATE INDEX idx_release_year_rating ON BUILDS (build_release_year, build_rating);
```

Findings and Explanation: The actual execution time decreased from 7.428 seconds to 6.136 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 3:

Index on Build\_Themes.theme\_id

Why: Since the Build\_Themes table is often joined with the BUILDS table, indexing theme\_id will speed up join operations when retrieving theme information related to builds.

```
CREATE INDEX idx_theme_id ON Build_Themes (theme_id);
```

Findings and Explanation: The actual execution time decreased from 5.288 seconds to 5.233 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

### **Query 3**

Default Index

- Cost: 1983.10

- Time: 74.770...74.775

```
mysql> EXPLAIN ANALYZE SELECT build_id, build_name, build_release_year, build_rating
-> FROM BUILDS
-> WHERE build_release_year >= 2019
-> AND build_rating > (
->   SELECT AVG(build_rating)
->   FROM BUILDS
->   WHERE build_release_year >= 2019
-> )
-> ORDER BY build_release_year DESC, build_rating DESC
-> LIMIT 15
->
-> ;

+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (cost=1983.10 rows=15) (actual time=74.770..74.775 rows=15 loops=1)
|   -> Sort: BUILDS.build_release_year DESC, BUILDS.build_rating DESC, limit input to 15 row(s) per chunk (cost=1983.10 rows=24362) (actual time=74.768..74.772
|     -> Filter: ((BUILDS.build_release_year >= 2019) and (BUILDS.build_rating > (select #2))) (cost=1983.10 rows=24362) (actual time=28.237..56.315 rows=3398
|       -> Table scan on BUILDS (cost=1983.10 rows=24362) (actual time=0.831..23.613 rows=24196 loops=1)
|         -> Select #2 (subquery in condition; run only once)
|           -> Aggregate: avg(BUILDS.build_rating) (cost=3336.44 rows=1) (actual time=24.226..24.228 rows=1 loops=1)
|             -> Filter: (BUILDS.build_release_year >= 2019) (cost=2524.45 rows=8120) (actual time=0.092..22.439 rows=6880 loops=1)
|               -> Table scan on BUILDS (cost=2524.45 rows=24362) (actual time=0.083..17.359 rows=24196 loops=1)
|
|
+-----+
(END)
```

Index 1:

Index on Build\_Details.part\_id and Build\_Details.build\_id

Why: This composite index will enhance the performance of queries that join the Build\_Details table with the BUILDS table, especially when retrieving part quantities and counting unique parts across builds.

CREATE INDEX idx\_part\_build ON Build\_Details (part\_id, build\_id);

Findings and Explanation: The actual execution time decreased from 5.288 seconds to 5.233 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

**Index 2:**

Index on BUILDS.build\_release\_year and BUILDS.build\_rating

Why: An index on these columns will speed up queries that filter on build\_release\_year and those that check for build\_rating, like the one that identifies builds released within a certain timeframe or above a certain rating.

CREATE INDEX idx\_release\_year\_rating ON BUILDS (build\_release\_year, build\_rating);

Findings and Explanation: The actual execution time decreased from 6.296 seconds to 4.126 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed

of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 3:

Index on Build\_Themes.theme\_id

Why: Since the Build\_Themes table is often joined with the BUILDS table, indexing theme\_id will speed up join operations when retrieving theme information related to builds.

```
CREATE INDEX idx_theme_id ON Build_Themes (theme_id);
```

Findings and Explanation: The actual execution time decreased from 5.288 seconds to 5.233 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

#### **Query 4**

Default Index

- Cost: 802.05
- Time: 7.836..7.840

Index 1:

Index on Build\_Details.part\_id and Build\_Details.build\_id

Why: This composite index will enhance the performance of queries that join the Build\_Details table with the BUILDS table, especially when retrieving part quantities and counting unique parts across builds.

```
CREATE INDEX idx_part_build ON Build_Details (part_id, build_id);
```

Findings and Explanation: The actual execution time decreased from 5.288 seconds to 5.233 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 2:

Index on BUILDS.build\_release\_year and BUILDS.build\_rating

Why: An index on these columns will speed up queries that filter on build\_release\_year and those that check for build\_rating, like the one that identifies builds released within a certain timeframe or above a certain rating.

```
CREATE INDEX idx_release_year_rating ON BUILDS (build_release_year,  
build_rating);
```

Findings and Explanation: The actual execution time decreased from 5.288 seconds to 5.233 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 3:

Index on Build\_Themes.theme\_id

Why: Since the Build\_Themes table is often joined with the BUILDS table, indexing theme\_id will speed up join operations when retrieving theme information related to builds.

```
CREATE INDEX idx_theme_id ON Build_Themes (theme_id);
```

Findings and Explanation: The actual execution time decreased from 6.296 seconds to 4.126 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

## Final Choice

After evaluating the various indexing configurations, we ultimately selected the index on the combined index on part\_id and build\_id as our final design. This decision was based on the consistent performance gains observed across multiple queries, particularly in those that part id and build id the most. The improvements were significant enough to justify the added overhead of maintaining the index, and we anticipate that it will enhance user experience by reducing query response times.