

DDL Commands

```
Create Table USERS (  
    user_id VARCHAR(100) Primary key,  
    username VARCHAR(100),  
    password VARCHAR(100),  
    user_email VARCHAR(100),  
    user_age INT  
);
```

```
Create table PARTS (  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    part_name VARCHAR(100),  
    part_png VARCHAR(255),  
    part_dimensions VARCHAR(255),  
    PRIMARY KEY(part_id, part_color)  
);
```

```
Create Table BUILDS (  
    build_id VARCHAR(100) Primary Key,  
    build_name VARCHAR(100),  
    build_png VARCHAR(100),  
    build_link VARCHAR(255),  
    build_age_rating INTEGER,  
    build_rating DECIMAL(4, 2),  
    build_release_year VARCHAR(10)  
);
```

```
Create Table THEMES (  
    theme_id VARCHAR(100) Primary Key,  
    theme_name VARCHAR(100),  
    theme_description VARCHAR(1000),  
    popular_build_id_1 VARCHAR(100),  
    popular_build_id_2 VARCHAR(100),  
    popular_build_id_3 VARCHAR(100),  
    FOREIGN KEY (popular_build_id_1) REFERENCES  
    BUILDS(build_id),
```

```
        FOREIGN KEY (popular_build_id_2) REFERENCES  
BUILDS(build_id),  
        FOREIGN KEY (popular_build_id_3) REFERENCES BUILDS(build_id)  
    );
```

```
Create Table SUPPLIERS (  
    supplier_id VARCHAR(100) Primary Key,  
    supplier_name VARCHAR(100),  
    supplier_region VARCHAR(100)  
);
```

```
Create Table INVENTORY (  
    user_id VARCHAR(100),  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    part_quantity INTEGER,  
    PRIMARY KEY(user_id, part_id, part_color),  
    FOREIGN KEY (user_id) REFERENCES USERS(user_id),  
    FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,  
part_color)  
);
```

```
Create Table BUILD_DETAILS (  
    build_id VARCHAR(100),  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    part_quantity INTEGER,  
    PRIMARY KEY(build_id, part_id, part_color),  
    FOREIGN KEY (build_id) REFERENCES BUILDS(build_id),  
    FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,  
part_color)  
);
```

```
Create Table FAVORITES(  
    user_id VARCHAR(100),  
    part_id VARCHAR(100),  
    part_color VARCHAR(100),  
    PRIMARY KEY(user_id, part_id, part_color),  
    FOREIGN KEY (user_id) REFERENCES USERS(user_id),
```

```
        FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,
part_color)
    );
    LOAD DATA LOCAL INFILE 'storage-legolab/data/PARTS.csv' INTO TABLE
table_name FIELDS TERMINATED BY ';' ENCLOSED BY '"' LINES
TERMINATED BY '\r\n'; IGNORE 1 ROWS;
```

```
Create Table REVIEWS(
    user_id VARCHAR(100),
    build_id VARCHAR(100),
    review_text VARCHAR(1000),
    PRIMARY KEY(user_id),
    FOREIGN KEY (user_id) REFERENCES USERS(user_id),
    FOREIGN KEY (build_id) REFERENCES BUILDS(build_id)
);
```

```
Create Table BUILD_PRICING (
    supplier_id VARCHAR(100),
    build_id VARCHAR(100),
    PRIMARY KEY(supplier_id, build_id),
    FOREIGN KEY (build_id) REFERENCES BUILDS(build_id),
    FOREIGN KEY (supplier_id) REFERENCES SUPPLIERS(supplier_id),
    build_price DECIMAL(6, 2)
);
```

```
Create Table PART_PRICING (
    supplier_id VARCHAR(100),
    part_id VARCHAR(100),
    part_color VARCHAR(100),
    PRIMARY KEY(supplier_id, part_id, part_color),
    FOREIGN KEY (part_id, part_color) REFERENCES PARTS(part_id,
part_color),
    FOREIGN KEY (supplier_id) REFERENCES SUPPLIERS(supplier_id),
    part_price DECIMAL(6, 2)
);
```

```
Create Table BUILD_HAS_THEME (
    theme_id VARCHAR(100),
    build_id VARCHAR(100),
    PRIMARY KEY(theme_id, build_id),
```

```
FOREIGN KEY (theme_id) REFERENCES THEMES(theme_id),
FOREIGN KEY (build_id) REFERENCES BUILDS(build_id)
);
```

GET the Fields of each Table

```
mysql> show tables;
+-----+
| Tables_in_legoLab |
+-----+
| BUILDS             |
| BUILD_DETAILS      |
| BUILD_HAS_THEME    |
| BUILD_PRICING      |
| FAVORITES          |
| INVENTORY          |
| PARTS              |
| PART_PRICING       |
| REVIEWS            |
| SUPPLIERS          |
| THEMES             |
| USERS              |
+-----+
12 rows in set (0.00 sec)
```

```
mysql> describe BUILD_DETAILS;
```

Field	Type	Null	Key	Default	Extra
user_id	varchar(100)	NO	PRI	NULL	
build_id	varchar(100)	NO	PRI	NULL	
part_quantity	int	YES		NULL	

```
3 rows in set (0.00 sec)
```

```
mysql> describe BUILD_HAS_THEME;
```

Field	Type	Null	Key	Default	Extra
theme_id	varchar(100)	NO	PRI	NULL	
build_id	varchar(100)	NO	PRI	NULL	

```
2 rows in set (0.00 sec)
```

```
mysql> describe BUILD_PRICING;
```

Field	Type	Null	Key	Default	Extra
supplier_id	varchar(100)	NO	PRI	NULL	
build_id	varchar(100)	NO	PRI	NULL	
build_price	decimal(6,2)	YES		NULL	

3 rows in set (0.01 sec)

```
mysql> describe FAVORITES;
```

Field	Type	Null	Key	Default	Extra
user_id	varchar(100)	NO	PRI	NULL	
part_id	varchar(100)	NO	PRI	NULL	
part_color	varchar(100)	NO	PRI	NULL	

3 rows in set (0.01 sec)

```
mysql> describe PARTS;
```

Field	Type	Null	Key	Default	Extra
part_id	varchar(100)	NO	PRI	NULL	
part_color	varchar(100)	NO	PRI	NULL	
part_name	varchar(100)	YES		NULL	
part_png	varchar(255)	YES		NULL	
part_dimensions	varchar(255)	YES		NULL	

5 rows in set (0.00 sec)

```
mysql> describe PART_PRICING;
```

Field	Type	Null	Key	Default	Extra
supplier_id	varchar(100)	NO	PRI	NULL	
part_id	varchar(100)	NO	PRI	NULL	
part_color	varchar(100)	NO	PRI	NULL	
part_price	decimal(6,2)	YES		NULL	

4 rows in set (0.00 sec)

```
mysql> describe REVIEWS;
```

Field	Type	Null	Key	Default	Extra
user_id	varchar(100)	NO	PRI	NULL	
build_id	varchar(100)	YES	MUL	NULL	
review_text	varchar(1000)	YES		NULL	

3 rows in set (0.01 sec)

```
mysql> describe SUPPLIERS;
```

Field	Type	Null	Key	Default	Extra
supplier_id	varchar(100)	NO	PRI	NULL	
supplier_name	varchar(100)	YES		NULL	
supplier_region	varchar(100)	YES		NULL	

3 rows in set (0.01 sec)

```
mysql> describe USERS;
```

Field	Type	Null	Key	Default	Extra
user_id	varchar(100)	NO	PRI	NULL	
username	varchar(100)	YES		NULL	
password	varchar(100)	YES		NULL	
user_email	varchar(100)	YES		NULL	
user_age	int	YES		NULL	

5 rows in set (0.01 sec)

```
mysql> describe THEMES;
```

Field	Type	Null	Key	Default	Extra
theme_id	varchar(100)	NO	PRI	NULL	
theme_name	varchar(100)	YES		NULL	
theme_description	varchar(1000)	YES		NULL	
popular_build_id_1	varchar(100)	YES	MUL	NULL	
popular_build_id_2	varchar(100)	YES	MUL	NULL	
popular_build_id_3	varchar(100)	YES	MUL	NULL	

6 rows in set (0.01 sec)

Count of Each Table

```
mysql> SELECT COUNT(*) FROM PARTS;
+-----+
| COUNT(*) |
+-----+
|      77949 |
+-----+
1 row in set (10.25 sec)
```

```
mysql> SELECT COUNT(*) FROM BUILDS;
+-----+
| COUNT(*) |
+-----+
|      24196 |
+-----+
1 row in set (4.79 sec)
```

```
mysql> SELECT COUNT(*) FROM THEMES;
+-----+
| COUNT(*) |
+-----+
|      1046 |
+-----+
1 row in set (0.01 sec)
```

```
mysql> SELECT COUNT(*) FROM USERS;
+-----+
| COUNT(*) |
+-----+
|          0 |
+-----+
1 row in set (0.08 sec)
```

```
mysql> SELECT COUNT(*) FROM SUPPLIERS;
+-----+
| COUNT(*) |
+-----+
|          0 |
+-----+
1 row in set (0.04 sec)
```

```
mysql> SELECT COUNT(*) FROM BUILD_HAS_THEME;
+-----+
| COUNT(*) |
+-----+
|      24198 |
+-----+
1 row in set (0.64 sec)
```

****Note** - We had to use small versions of these tables with 5000 to 10000 rows in order to actually get outputs within a decent amount of time and not have to run every single query for hours.

Queries:

1. Query: Find the Most Frequently Used Parts Across All Builds

This query identifies the parts most frequently used across all builds by joining the PARTS and Build_Details tables, aggregating by part_id and part_color.

SQL Concepts: **Join multiple relations, Aggregation with GROUP BY**

```
SELECT DISTINCT lp.part_id, lp.part_name, SUM(bd.part_quantity) AS
total_quantity_used
FROM small_PARTS AS lp
JOIN small_BUILD_DETAILS AS bd ON lp.part_id = bd.part_id
GROUP BY lp.part_id, lp.part_name
ORDER BY total_quantity_used DESC, lp.part_name
LIMIT 15;
```



```
mysql> SELECT DISTINCT lp.part_id, lp.part_name, SUM(bd.part_quantity) AS total_quantity_used
-> FROM small_PARTS AS lp
-> JOIN small_BUILD_DETAILS AS bd ON lp.part_id = bd.part_id
-> GROUP BY lp.part_id, lp.part_name
-> ORDER BY total_quantity_used DESC, lp.part_name
-> LIMIT 15;
```

part_id	part_name	total_quantity_used
2357	Brick 2 x 2 Corner	30118
2339	Brick Arch 1 x 5 x 4 [Continuous Bow]	870
2343	Equipment Goblet / Glass	870
2335	Flag 2 x 2 Square [Thin Clips]	728
2356	Brick 4 x 6	290
2340	Tail 4 x 1 x 3	128
132a	Tyre Smooth Old Style - Small	126
2345	Panel 3 x 3 x 6 Corner Wall	80
12825	Tile Special 1 x 1 with Clip with Rounded Tips	80
2039	Lamp Post 2 x 2 x 7 with 6 Base Flutes	63
22667	Plant, Cherries	40
2341	Slope Inverted 45° 3 x 1 Double with 2 Completely Open Studs	16
14226c11	String with End Studs 11L Overall	16
23306	Weapon Lightsaber Hilt without Bottom Ring	12
2342	Equipment Control Panel	8

15 rows in set (0.03 sec)

2. Query: Find Builds with the Largest Variety of Unique Parts

This query finds builds with the highest count of unique parts (based on part_id and part_color) used in each build, helping to identify builds that are the most complex in terms of part diversity.

SQL Concepts: **Join multiple relations, Aggregation with GROUP BY**

```
SELECT lb.build_id, lb.build_name, COUNT(DISTINCT bd.part_id, bd.part_color) AS
unique_part_count
FROM small_BUILDS AS lb
JOIN small_BUILD_DETAILS AS bd ON lb.build_id = bd.build_id
GROUP BY lb.build_id, lb.build_name
ORDER BY unique_part_count DESC, lb.build_name
LIMIT 15;
```

```
mysql> SELECT lb.build_id, lb.build_name, COUNT(DISTINCT bd.part_id, bd.part_color) AS unique_part_count
-> FROM small_BUILDS AS lb
-> JOIN small_BUILD_DETAILS AS bd ON lb.build_id = bd.build_id
-> GROUP BY lb.build_id, lb.build_name
-> ORDER BY unique_part_count DESC, lb.build_name
-> LIMIT 15;
```

build_id	build_name	unique_part_count
10184-1	Town Plan	402
10185-1	Green Grocer	297
10159-1	City Airport	293
10159-2	City Airport (Full Size Image Box)	293
10129-1	Rebel Snowspeeder	280
10179-1	Millennium Falcon	270
10019-1	Rebel Blockade Runner	247
10134-1	Y-wing Attack Starfighter	247
10144-1	Sandcrawler	220
10132-1	Motorized Hogwarts Express	208
10030-1	Imperial Star Destroyer	205
10123-1	Cloud City	184
10173-1	Holiday Train	183
10040-1	Black Seas Barracuda	179
10176-1	Royal King's Castle	178

15 rows in set (0.23 sec)

3. Query: Find Builds Released in the Last 5 Years with Above-Average Ratings

This query identifies builds released within the past five years that have a rating above the overall average for that time period.

SQL Concepts: **Subqueries, Aggregation with GROUP BY**

```
SELECT build_id, build_name, build_release_year, build_rating
FROM small_BUILDS
WHERE build_release_year >= 2019
AND build_rating > (
    SELECT AVG(build_rating)
    FROM small_BUILDS
    WHERE build_release_year >= 2019
)
ORDER BY build_release_year DESC, build_rating DESC
LIMIT 15;
```

```
mysql> SELECT build_id, build_name, build_release_year, build_rating FROM small_BUILDS WHERE build_release_year >= 2019 AND build_rating > (
ld_rating) FROM small_BUILDS WHERE build_release_year >= 2019 ) ORDER BY build_release_year DESC, build_rating DESC LIMIT 15;
```

build_id	build_name	build_release_year	build_rating
202292404-1	Ninjago Arin Backpack with Gym Bag and Pencil Case	2024	5.00
11038-1	Vibrant Creative Brick Box	2024	5.00
31147-1	Retro Camera	2024	5.00
10794-1	Team Spidey Web Spinner Headquarters	2024	5.00
202832409-1	Ninjago Red Duffle Bag	2024	4.90
30672-1	Steve and Baby Panda	2024	4.90
202792410-1	Friends Nova and Aliya Backpack	2024	4.80
21259-1	The Pirate Ship Voyage	2024	4.80
10795-1	Crafting with Baby Box	2024	4.80
242403-1	Thor	2024	4.70
11036-1	Creative Vehicles	2024	4.60
30665-1	Baby Gorilla Encounter	2024	4.60
2000409-2	Window Exploration Bag	2024	4.60
10332-1	Medieval Town Square	2024	4.50
122403-1	Owen with Helicopter	2024	4.50

```
15 rows in set (0.01 sec)
```

4. Find Builds with the Highest Number of Parts Used

This query lists the builds with the most parts used, based on the sum of part_quantity in Build_Details. This can be useful for identifying complex builds.

SQL Concepts: **Join multiple relations, Aggregation with GROUP BY**

```
SELECT lb.build_id, lb.build_name, SUM(bd.part_quantity) AS total_parts_used
FROM small_BUILDS AS lb
JOIN small_BUILD_DETAILS AS bd ON lb.build_id = bd.build_id
GROUP BY lb.build_id, lb.build_name
ORDER BY total_parts_used DESC
LIMIT 15;
```

```
mysql> SELECT lb.build_id, lb.build_name, SUM(bd.part_quantity) AS total_parts_used
-> FROM small_BUILDS AS lb
-> JOIN small_BUILD DETAILS AS bd ON lb.build_id = bd.build_id
-> GROUP BY lb.build_id, lb.build_name
-> ORDER BY total_parts_used DESC
-> LIMIT 15;
```

build_id	build_name	total_parts_used
10179-1	Millennium Falcon	4983
10143-1	Death Star II	3461
10181-1	Eiffel Tower 1:300	2875
10030-1	Imperial Star Destroyer	2804
10185-1	Green Grocer	2302
10184-1	Town Plan	1910
10018-1	Darth Maul	1867
10019-1	Rebel Blockade Runner	1741
10129-1	Rebel Snowspeeder	1414
10134-1	Y-wing Attack Starfighter	1352
10144-1	Sandcrawler	1322
10175-1	Vader's TIE Advanced	1212
10177-1	Boeing 787 Dreamliner	1102
10183-1	Hobby Train	1070
10178-1	Motorized Walking AT-AT	991

15 rows in set (0.03 sec)

INDEXING:

Query 1

Default Index

- Cost: 1005.25 for table scan
- Time: 5.65 on table scan

```
mysql> EXPLAIN ANALYZE SELECT DISTINCT lp.part_id, lp.part_name, SUM(bd.part_quantity) AS total_quantity_used FROM small PARTS AS lp JOIN small_BUILD DETAILS A
S bd ON lp.part_id = bd.part_id GROUP BY lp.part_id, lp.part_name ORDER BY total_quantity_used DESC, lp.part_name LIMIT 15;
```

EXPLAIN
-> Limit: 15 row(s) (actual time=20.836..20.839 rows=15 loops=1)
-> Sort: total_quantity_used DESC, lp.part_name, limit input to 15 row(s) per chunk (actual time=20.835..20.836 rows=15 loops=1)
-> Table scan on <temporary> (actual time=20.746..20.755 rows=34 loops=1)
-> Aggregate using temporary table (actual time=20.743..20.743 rows=34 loops=1)
-> Filter: (lp.part_id = bd.part_id) (cost=9660584.54 rows=9658926) (actual time=7.753..15.880 rows=3225 loops=1)
-> Inner hash join (<hash>(lp.part_id)=<hash>(bd.part_id)) (cost=9660584.54 rows=9658926) (actual time=7.749..15.256 rows=3225 loops=1)
-> Table scan on lp (cost=0.08 rows=9846) (actual time=0.023..5.931 rows=10000 loops=1)
-> Hash
-> Table scan on bd (cost=1005.25 rows=9810) (actual time=0.037..5.650 rows=10000 loops=1)

1 row in set (0.02 sec)

Index 1:

Index on Build_Details.part_id

Why: This index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when retrieving part quantities.

CREATE INDEX idx_part_details ON small_BUILD_DETAILS (part_id);

```
mysql> EXPLAIN ANALYZE SELECT DISTINCT lp.part_id, lp.part_name, SUM(bd.part_quantity) AS total_quantity_used FROM small PARTS AS lp JOIN small_BUILD_DETAILS AS bd ON lp.part_id = bd.part_id GROUP BY lp.part_id, lp.part_name ORDER BY total_quantity_used DESC, lp.part_name LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (actual time=44.016..44.019 rows=15 loops=1) |
|   -> Sort: total_quantity_used DESC, lp.part_name, limit input to 15 row(s) per chunk (actual time=44.016..44.017 rows=15 loops=1) |
|     -> Table scan on <temporary> (actual time=43.965..43.975 rows=34 loops=1) |
|       -> Aggregate using temporary table (actual time=43.962..43.962 rows=34 loops=1) |
|         -> Nested loop inner join (cost=25522.13 rows=69992) (actual time=1.395..38.322 rows=3225 loops=1) |
|           -> Table scan on lp (cost=1024.85 rows=9846) (actual time=0.058..6.130 rows=10000 loops=1) |
|             -> Index lookup on bd using idx_part_details (part_id=lp.part_id) (cost=1.78 rows=7) (actual time=0.002..0.003 rows=0 loops=10000) |
|   |
+-----+
| (END) |
+-----+
```

Findings and Explanation: The actual execution time decreased from 5.65 seconds to 0.002 seconds, showing a measurable improvement. The query cost decreased considerably to 1.78 from 1005. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 2:

Index on Build_Details.build_id

Why: This index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when counting unique parts across builds

CREATE INDEX idx_build_details ON Build_Details (build_id);

```
mysql> EXPLAIN ANALYZE SELECT DISTINCT lp.part_id, lp.part_name, SUM(bd.part_quantity) AS total_quantity_used FROM small PARTS AS lp JOIN small_BUILD_DETAILS AS bd ON lp.part_id = bd.part_id GROUP BY lp.part_id, lp.part_name ORDER BY total_quantity_used DESC, lp.part_name LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (actual time=21.624..21.627 rows=15 loops=1) |
|   -> Sort: total_quantity_used DESC, lp.part_name, limit input to 15 row(s) per chunk (actual time=21.624..21.625 rows=15 loops=1) |
|     -> Table scan on <temporary> (actual time=21.576..21.583 rows=34 loops=1) |
|       -> Aggregate using temporary table (actual time=21.573..21.573 rows=34 loops=1) |
|         -> Filter: (lp.part_id = bd.part_id) (cost=9660584.54 rows=9658926) (actual time=8.609..16.771 rows=3225 loops=1) |
|           -> Inner hash join (<hash>(lp.part_id)=<hash>(bd.part_id)) (cost=9660584.54 rows=9658926) (actual time=8.605..16.173 rows=3225 loops=1) |
|             -> Table scan on lp (cost=0.08 rows=9846) (actual time=0.021..5.945 rows=10000 loops=1) |
|               -> Hash |
|             -> Table scan on bd (cost=1005.25 rows=9810) (actual time=0.038..5.923 rows=10000 loops=1) |
|   |
+-----+
| (END) |
+-----+
```

Findings and Explanation: The actual execution time increased from 5.65 seconds to 5.78 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy is not that great compared to the default indexing

Index 3:

Index on Build_Details.part_id and Build_Details.build_id

Why: This composite index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when retrieving part quantities and counting unique parts across builds

CREATE INDEX idx_part_build_details ON small_BUILD_DETAILS(part_id, build_id);

```
mysql> EXPLAIN ANALYZE SELECT DISTINCT lp.part_id, lp.part_name, SUM(bd.part_quantity) AS total_quantity_used FROM small PARTS AS lp JOIN small_BUILD_DETAILS AS bd ON lp.part_id = bd.part_id GROUP BY lp.part_id, lp.part_name ORDER BY total_quantity_used DESC, lp.part_name LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (actual time=41.054..41.057 rows=15 loops=1)
  -> Sort: total_quantity_used DESC, lp.part_name, limit input to 15 row(s) per chunk (actual time=41.053..41.055 rows=15 loops=1)
    -> Table scan on <temporary> (actual time=40.988..41.001 rows=34 loops=1)
      -> Aggregate using temporary table (actual time=40.983..40.983 rows=34 loops=1)
        -> Nested loop inner join (cost=25522.13 rows=69992) (actual time=0.443..35.535 rows=3225 loops=1)
          -> Table scan on lp (cost=1024.85 rows=9846) (actual time=0.043..6.253 rows=10000 loops=1)
            -> Index lookup on bd using idx_part_build_details (part_id=lp.part_id) (cost=1.78 rows=7) (actual time=0.002..0.003 rows=0 loops=10000)
          |
        |
      |
    |
  |
+-----+
| (END) |
+-----+
```

Findings and Explanation: The actual execution time decreased from 5.65 seconds to 0.03 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Query 2

Default Index

- Cost: 504.95 on scan
- Time: 44.07 on limit

```
mysql> EXPLAIN ANALYZE SELECT lb.build_id, lb.build_name, COUNT(DISTINCT bd.part_id, bd.part_color) AS unique_part_count FROM small_BUILDS AS lb JOIN small_BUILD_DETAILS AS bd ON lb.build_id = bd.build_id GROUP BY lb.build_id, lb.build_name ORDER BY unique_part_count DESC, lb.build_name LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (actual time=44.067..44.070 rows=15 loops=1)
  -> Sort: unique_part_count DESC, lb.build_name, limit input to 15 row(s) per chunk (actual time=44.066..44.069 rows=15 loops=1)
    -> Stream results (actual time=30.724..43.914 rows=158 loops=1)
      -> Group aggregate: count(distinct small_BUILD_DETAILS.part_id,small_BUILD_DETAILS.part_color) (actual time=30.720..43.853 rows=158 loops=1)
        -> Sort: lb.build_id, lb.build_name (actual time=30.703..31.829 rows=10000 loops=1)
          -> Stream results (cost=4716554.17 rows=4715667) (actual time=4.858..24.330 rows=10000 loops=1)
            -> Filter: (bd.build_id = lb.build_id) (cost=4716554.17 rows=4715667) (actual time=4.851..19.438 rows=10000 loops=1)
              -> Inner hash join (<hash>(bd.build_id)=<hash>(lb.build_id)) (cost=4716554.17 rows=4715667) (actual time=4.848..17.014 rows=10000 loops=1)
                -> Table scan on bd (cost=0.10 rows=9810) (actual time=0.019..6.137 rows=10000 loops=1)
                  -> Hash
                -> Table scan on lb (cost=504.95 rows=4807) (actual time=0.032..3.205 rows=5000 loops=1)
              |
            |
          |
        |
      |
    |
  |
+-----+
| (END) |
+-----+
```

Index 1:

Index on Build_Details.part_id

Why: This index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when retrieving part quantities.

CREATE INDEX idx_part_details ON Build_Details (part_id);

```
mysql> EXPLAIN ANALYZE SELECT lb.build_id, lb.build_name, COUNT(DISTINCT bd.part_id, bd.part_color) AS unique_part_count FROM small_BUILDS AS lb JOIN small_BUILD_DETAILS AS bd ON lb.build_id = bd.build_id GROUP BY lb.build_id, lb.build_name ORDER BY unique_part_count DESC, lb.build_name LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (actual time=44.602..44.606 rows=15 loops=1)
  -> Sort: unique_part_count DESC, lb.build_name, limit input to 15 row(s) per chunk (actual time=44.601..44.604 rows=15 loops=1)
    -> Stream results (actual time=30.957..44.501 rows=158 loops=1)
      -> Group aggregate: count(distinct small_BUILD_DETAILS.part_id,small_BUILD_DETAILS.part_color) (actual time=30.953..44.385 rows=158 loops=1)
        -> Sort: lb.build_id, lb.build_name (actual time=30.934..32.060 rows=10000 loops=1)
          -> Stream results (cost=4716554.17 rows=4715667) (actual time=4.604..24.320 rows=10000 loops=1)
            -> Filter: (bd.build_id = lb.build_id) (cost=4716554.17 rows=4715667) (actual time=4.596..19.425 rows=10000 loops=1)
              -> Inner hash join (<hash>(bd.build_id)=<hash>(lb.build_id)) (cost=4716554.17 rows=4715667) (actual time=4.592..17.077 rows=10000 loops=1)
                -> Table scan on bd (cost=0.10 rows=9810) (actual time=0.019..6.351 rows=10000 loops=1)
                  -> Hash
                -> Table scan on lb (cost=504.95 rows=4807) (actual time=0.029..2.970 rows=5000 loops=1)
              |
            |
          |
        |
      |
    |
  |
+-----+
| (END) |
+-----+
```

Findings and Explanation: The actual execution time decreased from 44.07 seconds to 2.97 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 2:

Index on Build_Details.build_id

Why: This index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when counting unique parts across builds

CREATE INDEX idx_build_details ON Build_Details (build_id);

```
mysql> EXPLAIN ANALYZE SELECT lb.build_id, lb.build_name, COUNT(DISTINCT bd.part_id, bd.part_color) AS unique_part_count FROM small_BUILDS AS lb JOIN small_BUILD_DETAILS AS bd ON lb.build_id = bd.build_id GROUP BY lb.build_id, lb.build_name ORDER BY unique_part_count DESC, lb.build_name LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s)  (actual time=52.248..52.252 rows=15 loops=1) |
|   -> Sort: unique_part_count DESC, lb.build_name, limit input to 15 row(s) per chunk  (actual time=52.248..52.251 rows=15 loops=1) |
|     -> Stream results  (cost=134811.92 rows=298460) (actual time=9.641..52.125 rows=158 loops=1) |
|       -> Group aggregate: count(distinct bd.part_id, bd.part_color)  (cost=134811.92 rows=298460) (actual time=9.637..52.033 rows=158 loops=1) |
|         -> Nested loop inner join  (cost=104965.93 rows=298460) (actual time=9.613..39.124 rows=10000 loops=1) |
|           -> Sort: lb.build_id, lb.build_name  (cost=504.95 rows=4807) (actual time=9.538..10.069 rows=5000 loops=1) |
|             -> Table scan on lb  (cost=504.95 rows=4807) (actual time=0.032..3.609 rows=5000 loops=1) |
|               -> Index lookup on bd using idx_build_details (build_id=lb.build_id)  (cost=15.52 rows=62) (actual time=0.003..0.006 rows=2 loops=5000) |
+-----+
| (END) |
+-----+
```

Findings and Explanation: The actual execution time decreased from 44.07 seconds to 3.609 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Index 3:

Index on Build_Details.part_id and Build_Details.build_id

Why: This composite index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when retrieving part quantities and counting unique parts across builds

CREATE INDEX idx_part_build_details ON Build_Details (part_id, build_id);

```
mysql> EXPLAIN ANALYZE SELECT lb.build_id, lb.build_name, COUNT(DISTINCT bd.part_id, bd.part_color) AS unique_part_count FROM small_BUILDS AS lb JOIN small_BUILD_DETAILS AS bd ON lb.build_id = bd.build_id GROUP BY lb.build_id, lb.build_name ORDER BY unique_part_count DESC, lb.build_name LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (actual time=43.947..43.950 rows=15 loops=1) |
|   -> Sort: unique part count DESC, lb.build_name, limit input to 15 row(s) per chunk (actual time=43.947..43.949 rows=15 loops=1) |
|     -> Stream results (actual time=30.586..43.869 rows=158 loops=1) |
|       -> Group aggregate: count(distinct small_BUILD_DETAILS.part_id, small_BUILD_DETAILS.part_color) (actual time=30.582..43.804 rows=158 loops=1) |
|         -> Sort: lb.build_id, lb.build_name (actual time=30.561..31.671 rows=10000 loops=1) |
|           -> Stream results (cost=4716554.17 rows=4715667) (actual time=4.472..24.333 rows=10000 loops=1) |
|             -> Filter: (bd.build_id = lb.build_id) (cost=4716554.17 rows=4715667) (actual time=4.466..19.266 rows=10000 loops=1) |
|               -> Inner hash join (<hash>(bd.build_id)=<hash>(lb.build_id)) (cost=4716554.17 rows=4715667) (actual time=4.463..16.914 rows=10000 loops=1) |
|                 -> Table scan on bd (cost=0.10 rows=9810) (actual time=0.017..6.265 rows=10000 loops=1) |
|                   -> Hash |
|                     -> Table scan on lb (cost=504.95 rows=4807) (actual time=0.034..2.994 rows=5000 loops=1) |
|   | |
+-----+
| (END) |
+-----+
```

Findings and Explanation: The actual execution time decreased from 44.07 seconds to 2.994 seconds, showing a measurable improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Query 3

Default Index

- Cost: 398.13 on limit
- Time: 7.458 on limit

```
mysql> EXPLAIN ANALYZE SELECT build_id, build_name, build_release_year, build_rating FROM small_BUILDS WHERE build_release_year >= 2019 AND build_rating > (SELECT AVG(build_rating) FROM small_BUILDS WHERE build_release_year >= 2019) ORDER BY build_release_year DESC, build_rating DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (cost=398.13 rows=15) (actual time=7.450..7.458 rows=15 loops=1) |
|   -> Sort: small_BUILDS.build_release_year DESC, small_BUILDS.build_rating DESC, limit input to 15 row(s) per chunk (cost=398.13 rows=4807) (actual time=7.447..7.450 rows=15 loops=1) |
|     -> Filter: ((small_BUILDS.build_release_year >= 2019) and (small_BUILDS.build_rating > (select #2))) (cost=398.13 rows=4807) (actual time=3.610..7.241 rows=15 loops=1) |
|       -> Table scan on small_BUILDS (cost=398.13 rows=4807) (actual time=0.037..3.039 rows=5000 loops=1) |
|         -> Select #2 (subquery in condition; run only once) |
|           -> Aggregate: avg(small_BUILDS.build_rating) (cost=665.17 rows=1) (actual time=3.525..3.526 rows=1 loops=1) |
|             -> Filter: (small_BUILDS.build_release_year >= 2019) (cost=504.95 rows=1602) (actual time=0.006..3.344 rows=1082 loops=1) |
|               -> Table scan on small_BUILDS (cost=504.95 rows=4807) (actual time=0.004..2.862 rows=5000 loops=1) |
|   | |
+-----+
| (END) |
+-----+
```

Index 1:

Index on Builds.builds_release_year

Why: This index will help improve the speed of the query when we are checking by indexing through build_release_year.

CREATE INDEX idx_build_year_ON small_BUILDS(build_release_year);

```
mysql> EXPLAIN ANALYZE SELECT build_id, build_name, build_release_year, build_rating FROM small_BUILDS WHERE build_release_year >= 2019 AND build_rating > (SELECT AVG(build_rating) FROM small_BUILDS WHERE build_release_year >= 2019) ORDER BY build_release_year DESC, build_rating DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (cost=398.13 rows=15) (actual time=7.404..7.407 rows=15 loops=1) |
|   -> Sort: small_BUILDS.build_release_year DESC, small_BUILDS.build_rating DESC, limit input to 15 row(s) per chunk (cost=398.13 rows=4807) (actual time=7.401..7.404 rows=15 loops=1) |
|     -> Filter: ((small_BUILDS.build_release_year >= 2019) and (small_BUILDS.build_rating > (select #2))) (cost=398.13 rows=4807) (actual time=3.654..7.201 rows=15 loops=1) |
|       -> Table scan on small_BUILDS (cost=398.13 rows=4807) (actual time=0.081..3.015 rows=5000 loops=1) |
|         -> Select #2 (subquery in condition; run only once) |
|           -> Aggregate: avg(small_BUILDS.build_rating) (cost=665.17 rows=1) (actual time=3.524..3.525 rows=1 loops=1) |
|             -> Filter: (small_BUILDS.build_release_year >= 2019) (cost=504.95 rows=1602) (actual time=0.006..3.386 rows=1082 loops=1) |
|               -> Table scan on small_BUILDS (cost=504.95 rows=4807) (actual time=0.004..2.857 rows=5000 loops=1) |
|   | |
+-----+
1 row in set, 2 warnings (0.01 sec)
```


Findings and Explanation: The actual execution time did not decrease at all. The query cost remained roughly the same. The Indexing strategy did not cause a significant improvement.

Index 2:

Index on `BUILDS.build_rating`

Why: An index on these columns will speed up queries that filter on `build_rating`, like the one that identifies builds released above a certain rating.

`CREATE INDEX idx_build_rating ON small_BUILDS(build_rating);`

```
mysql> EXPLAIN ANALYZE SELECT build_id, build_name, build_release_year, build_rating FROM small_BUILDS WHERE build_release_year >= 2019 AND build_rating > (
  SELECT AVG(build_rating) FROM small_BUILDS WHERE build_release_year >= 2019 ) ORDER BY build_release_year DESC, build_rating DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (cost=504.95 rows=15) (actual time=4.017..4.019 rows=15 loops=1)
  -> Sort: small_BUILDS.build_release_year DESC, small_BUILDS.build_rating DESC, limit input to 15 row(s) per chunk (cost=504.95 rows=4807) (actual time=4.017..4.019 rows=15 loops=1)
  -> Filter: ((small_BUILDS.build_release_year >= 2019) and (small_BUILDS.build_rating > (select #2))) (cost=504.95 rows=4807) (actual time=0.054..3.811 rows=15 loops=1)
  -> Table scan on small_BUILDS (cost=504.95 rows=4807) (actual time=0.024..3.185 rows=5000 loops=1)
  -> Select #2 (subquery in condition; run only once)
    -> Aggregate: avg(small_BUILDS.build_rating) (cost=665.17 rows=1) (actual time=3.586..3.586 rows=1 loops=1)
    -> Filter: (small_BUILDS.build_release_year >= 2019) (cost=504.95 rows=1602) (actual time=0.049..3.449 rows=1082 loops=1)
    -> Table scan on small_BUILDS (cost=504.95 rows=4807) (actual time=0.037..2.886 rows=5000 loops=1)
+-----+
1 row in set, 2 warnings (0.01 sec)
```

Findings and Explanation: The actual execution time remained the same. The query cost actually increased. The indexing strategy reduced the speed of the query by decreasing the speed of joins, aggregations, and sorting operations, leading to slower execution times and less efficient resource usage.

Index 3:

Index on `BUILDS.builds_release_year` and `BUILDS.builds_rating`

Why: Since the `Build_Themes` table is often joined with the `BUILDS` table, indexing `theme_id` will speed up join operations when retrieving theme information related to builds.

`CREATE INDEX idx_release_year_rating ON BUILDS (build_release_year, build_rating);`

```
mysql> CREATE INDEX idx_release_year_rating ON small_BUILDS (build_release_year, build_rating);
Query OK, 0 rows affected (0.17 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> EXPLAIN ANALYZE SELECT build_id, build_name, build_release_year, build_rating FROM small_BUILDS WHERE build_release_year >= 2019 AND build_rating > (
  SELECT AVG(build_rating) FROM small_BUILDS WHERE build_release_year >= 2019 ) ORDER BY build_release_year DESC, build_rating DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (cost=2.56 rows=5) (actual time=0.060..0.133 rows=15 loops=1)
  -> Filter: ((small_BUILDS.build_release_year >= 2019) and (small_BUILDS.build_rating > (select #2))) (cost=2.56 rows=5) (actual time=0.059..0.130 rows=15 loops=1)
  -> Index scan on small_BUILDS using idx_release_year_rating (reverse) (cost=2.56 rows=29) (actual time=0.050..0.115 rows=15 loops=1)
  -> Select #2 (subquery in condition; run only once)
    -> Aggregate: avg(small_BUILDS.build_rating) (cost=665.17 rows=1) (actual time=5.381..5.381 rows=1 loops=1)
    -> Filter: (small_BUILDS.build_release_year >= 2019) (cost=504.95 rows=1602) (actual time=4.033..5.151 rows=1082 loops=1)
    -> Covering index scan on small_BUILDS using idx_release_year_rating (cost=504.95 rows=4807) (actual time=0.036..4.216 rows=5000 loops=1)
+-----+
1 row in set, 0 warnings (0.01 sec)
```

Findings and Explanation: The actual execution time decreased greatly from 7.45 seconds to 0.06 seconds, showing a measurable improvement. The query cost also decreased considerably to 2.56 from 398.13. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Query 4

Default Index

- Cost: 504.95 on table scan
- Time: 29.573 on limit

```
mysql> EXPLAIN ANALYZE SELECT lb.build_id, lb.build_name, SUM(bd.part_quantity) AS total_parts_used FROM small_BUILDS AS lb JOIN small_BUILD_DETAILS AS bd ON lb.build_id = bd.build_id GROUP BY lb.build_id, lb.build_name ORDER BY total_parts_used DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (actual time=29.571..29.573 rows=15 loops=1) |
|   -> Sort: total_parts_used DESC, limit input to 15 row(s) per chunk (actual time=29.570..29.571 rows=15 loops=1) |
|     -> Table scan on <temporary> (actual time=29.450..29.510 rows=158 loops=1) |
|       -> Aggregate using temporary table (actual time=29.446..29.446 rows=158 loops=1) |
|         -> Filter: (bd.build_id = lb.build_id) (cost=4716554.17 rows=4715667) (actual time=4.748..18.837 rows=10000 loops=1) |
|           -> Inner hash join (<hash>(bd.build_id)=<hash>(lb.build_id)) (cost=4716554.17 rows=4715667) (actual time=4.745..16.389 rows=10000 loops=1) |
|             -> Table scan on bd (cost=0.10 rows=9810) (actual time=0.020..5.925 rows=10000 loops=1) |
|               -> Hash |
|             -> Table scan on lb (cost=504.95 rows=4807) (actual time=0.031..3.172 rows=5000 loops=1) |
|   |
+-----+
(END)
```

Index 1:

Index on Build_Details.part_id

Why: This index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when retrieving part quantities.

CREATE INDEX idx_part_details ON Build_Details (part_id);

```
mysql> EXPLAIN ANALYZE SELECT lb.build_id, lb.build_name, SUM(bd.part_quantity) AS total_parts_used FROM small_BUILDS AS lb JOIN small_BUILD_DETAILS AS bd ON lb.build_id = bd.build_id GROUP BY lb.build_id, lb.build_name ORDER BY total_parts_used DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (actual time=36.705..36.709 rows=15 loops=1) |
|   -> Sort: total_parts_used DESC, limit input to 15 row(s) per chunk (actual time=36.704..36.706 rows=15 loops=1) |
|     -> Table scan on <temporary> (actual time=36.507..36.594 rows=158 loops=1) |
|       -> Aggregate using temporary table (actual time=36.479..36.479 rows=158 loops=1) |
|         -> Filter: (bd.build_id = lb.build_id) (cost=4716554.17 rows=4715667) (actual time=4.610..23.706 rows=10000 loops=1) |
|           -> Inner hash join (<hash>(bd.build_id)=<hash>(lb.build_id)) (cost=4716554.17 rows=4715667) (actual time=4.606..20.074 rows=10000 loops=1) |
|             -> Table scan on bd (cost=0.10 rows=9810) (actual time=0.018..6.959 rows=10000 loops=1) |
|               -> Hash |
|             -> Table scan on lb (cost=504.95 rows=4807) (actual time=0.028..3.071 rows=5000 loops=1) |
|   |
+-----+
(END)
```

Findings and Explanation: The actual execution time decreased from 29.573 seconds to 36.709 seconds, showing increased runtime. The query cost remained roughly the same. The indexing strategy was not optimized and the query decreased in speed with less efficient resource use.

Index 2:

Index on Build_Details.build_id

Why: This index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when counting unique parts across builds

CREATE INDEX idx_build_details ON Build_Details (build_id);

```
mysql> EXPLAIN ANALYZE SELECT lb.build_id, lb.build_name, SUM(bd.part_quantity) AS total_parts_used FROM small_BUILDS AS lb JOIN small_BUILD_DETAILS AS bd ON lb.build_id = bd.build_id GROUP BY lb.build_id, lb.build_name ORDER BY total_parts_used DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (actual time=45.079..45.081 rows=15 loops=1) |
|   -> Sort: total_parts_used DESC, limit input to 15 row(s) per chunk (actual time=45.078..45.079 rows=15 loops=1) |
|     -> Table scan on <temporary> (actual time=44.969..45.019 rows=158 loops=1) |
|       -> Aggregate using temporary table (actual time=44.964..44.964 rows=158 loops=1) |
|         -> Nested loop inner join (cost=104965.93 rows=298460) (actual time=0.082..33.694 rows=10000 loops=1) |
|           -> Table scan on lb (cost=504.95 rows=4807) (actual time=0.030..3.374 rows=5000 loops=1) |
|             -> Index lookup on bd using idx_build_details (build_id=lb.build_id) (cost=15.52 rows=62) (actual time=0.003..0.006 rows=2 loops=5000) |
|   |
+-----+
| (END) |
+-----+
```

Findings and Explanation: The actual execution time decreased from 29.573 seconds to 45.079 seconds, showing increased runtime. The query cost remained roughly the same. The indexing strategy was not optimized and the query decreased in speed with less efficient resource use.

Index 3:

Index on Build_Details.part_id and Build_Details.build_id

Why: This composite index will enhance the performance of queries that join the Build_Details table with the BUILDS table, especially when retrieving part quantities and counting unique parts across builds

CREATE INDEX idx_part_build_details ON Build_Details (part_id, build_id);

```
mysql> EXPLAIN ANALYZE SELECT lb.build_id, lb.build_name, SUM(bd.part_quantity) AS total_parts_used FROM small_BUILDS AS lb JOIN small_BUILD_DETAILS AS bd ON lb.build_id = bd.build_id GROUP BY lb.build_id, lb.build_name ORDER BY total_parts_used DESC LIMIT 15;
+-----+
| EXPLAIN |
+-----+
| -> Limit: 15 row(s) (actual time=29.361..29.363 rows=15 loops=1) |
|   -> Sort: total_parts_used DESC, limit input to 15 row(s) per chunk (actual time=29.360..29.361 rows=15 loops=1) |
|     -> Table scan on <temporary> (actual time=29.259..29.306 rows=158 loops=1) |
|       -> Aggregate using temporary table (actual time=29.255..29.255 rows=158 loops=1) |
|         -> Filter: (bd.build_id = lb.build_id) (cost=4716554.17 rows=4715667) (actual time=4.666..18.607 rows=10000 loops=1) |
|           -> Inner hash join (<hash>(bd.build_id)=<hash>(lb.build_id)) (cost=4716554.17 rows=4715667) (actual time=4.661..16.231 rows=10000 loops=1) |
|             -> Table scan on bd (cost=0.10 rows=9810) (actual time=0.031..5.942 rows=10000 loops=1) |
|               -> Hash |
|             -> Table scan on lb (cost=504.95 rows=4807) (actual time=0.029..3.038 rows=5000 loops=1) |
|   |
+-----+
| (END) |
+-----+
```

Findings and Explanation: The actual execution time decreased from 29.573 seconds to 29.363 seconds, showing a slight improvement. The query cost remained roughly the same. The indexing strategy effectively optimized the query by improving the speed of joins, aggregations, and sorting operations, leading to faster execution times and more efficient resource usage.

Final Choice

After evaluating the various indexing configurations, we ultimately selected the index on the combined index on part_id and build_id for build_details for BUILD_DETAILS table and combined index on build_year_release and build_rating for BUILDS table as our final design. This decision was based on the consistent performance gains observed across multiple queries, particularly in those that build_release_year and build_rating the most. The improvements were significant enough to justify the added overhead of maintaining the index, and we anticipate that it will enhance user experience by reducing query response times.